



Lecture: Pixels and Filters

Juan Carlos Niebles and Ranjay Krishna
Stanford Vision Lab



What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



Types of Images

Binary



Gray Scale

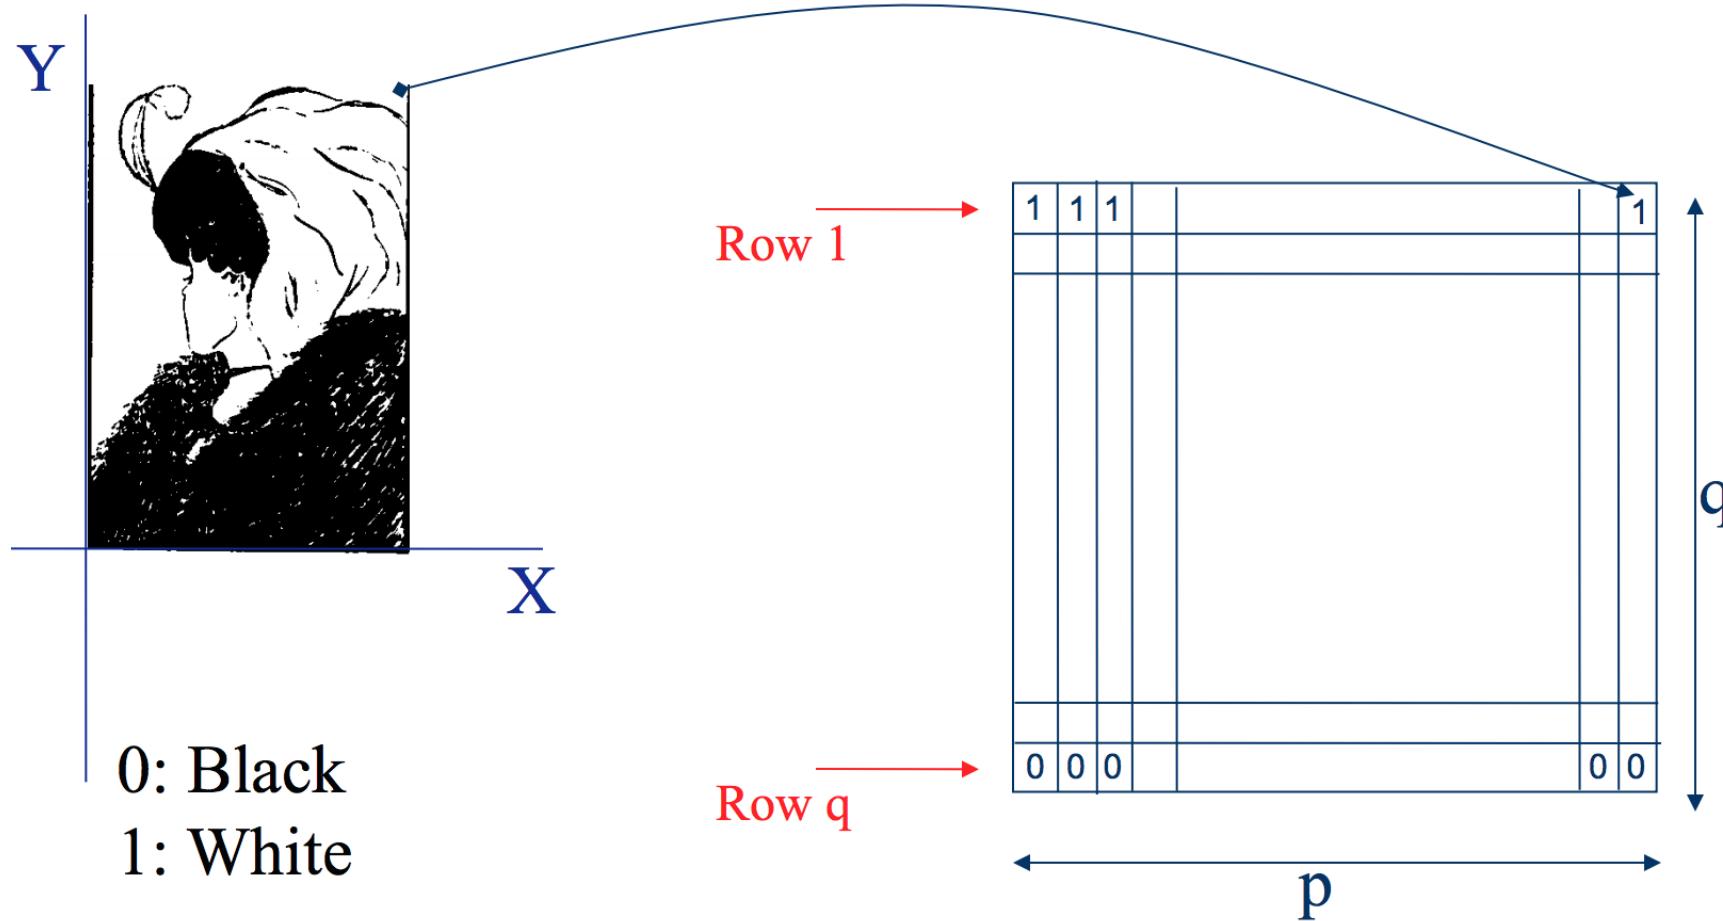


Color





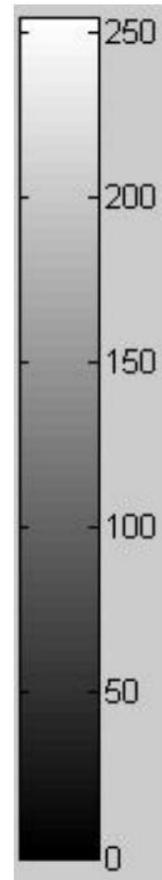
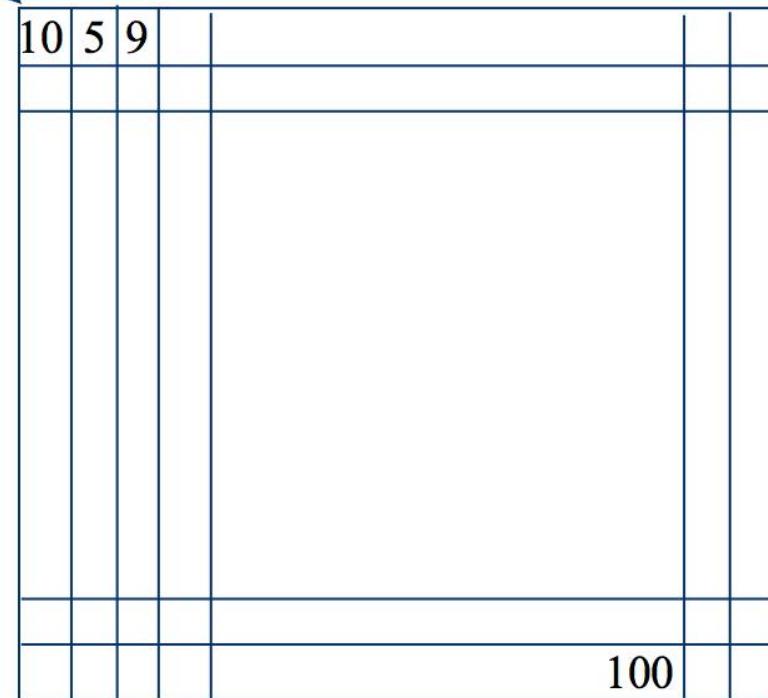
Binary image representation



Slide credit: Ulas Bagci



Grayscale image representation



Slide credit: Ulas Bagci



Color Image - one channel



Phil Noble / AP





Color image representation



Phil Noble / AP





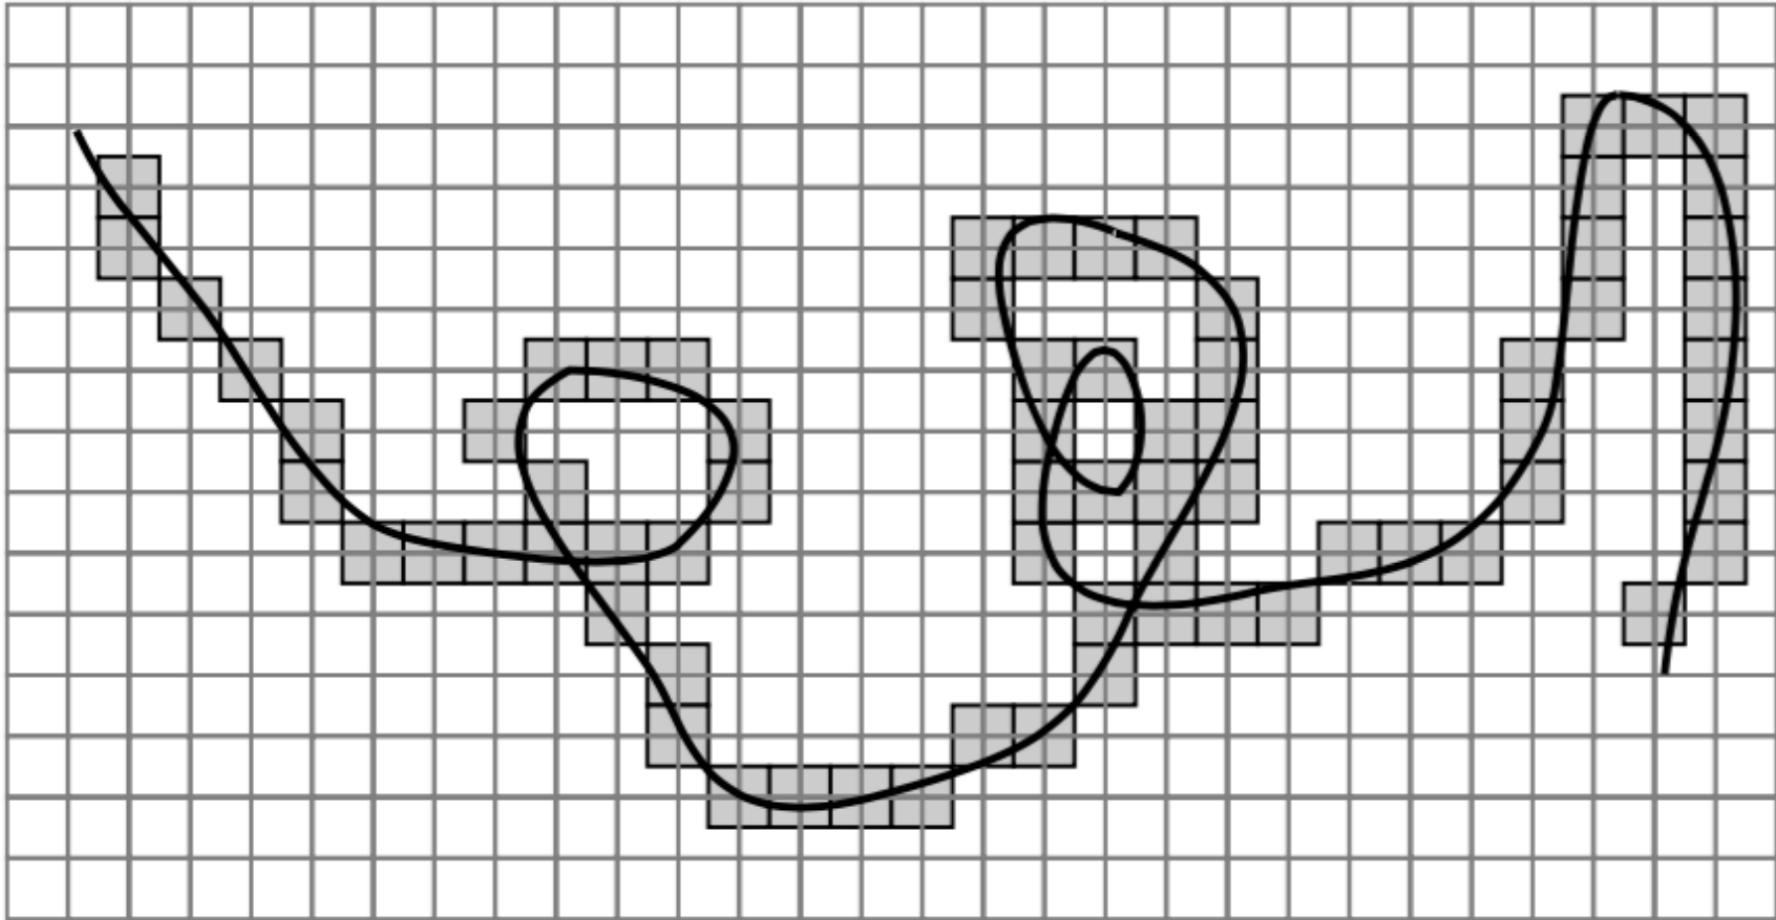
Images are sampled

What happens when we zoom
into the images we capture?





Errors due Sampling





Resolution

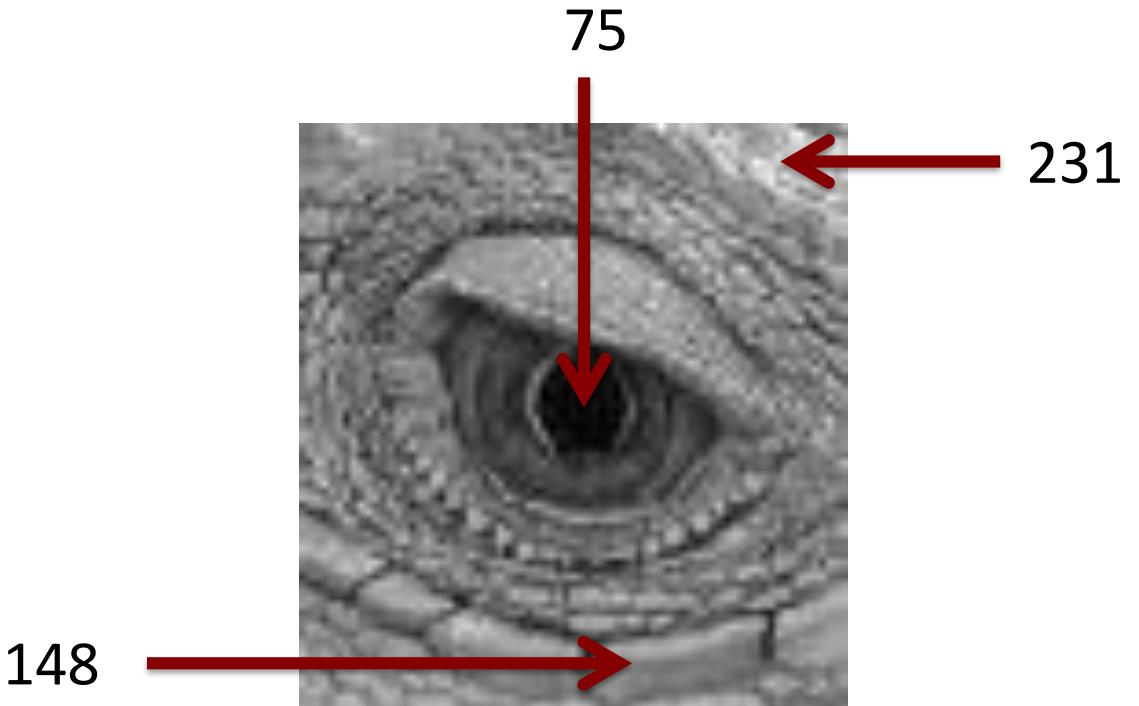
is a **sampling** parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density, and its standard value for recent screen technologies is 72 dpi





Images are Sampled and Quantized

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
(or “intensity”): [0,255]





Images are Sampled and Quantized

- An image contains discrete number of pixels

- A simple example

- Pixel value:

- “grayscale”

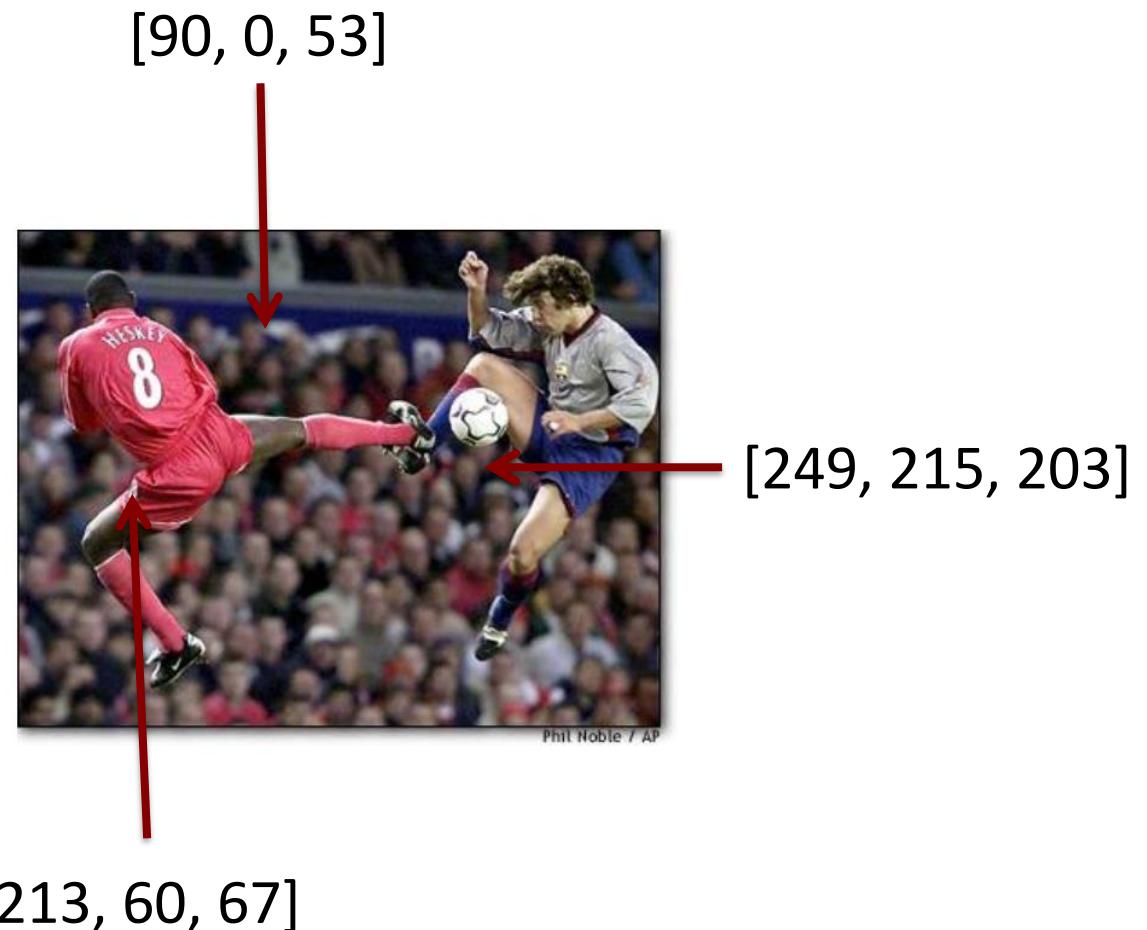
- (or “intensity”): [0,255]

- “color”

- RGB: [R, G, B]

- Lab: [L, a, b]

- HSV: [H, S, V]





With this loss of information (from sampling and quantization),

Can we still use images for useful tasks?



What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



Histogram

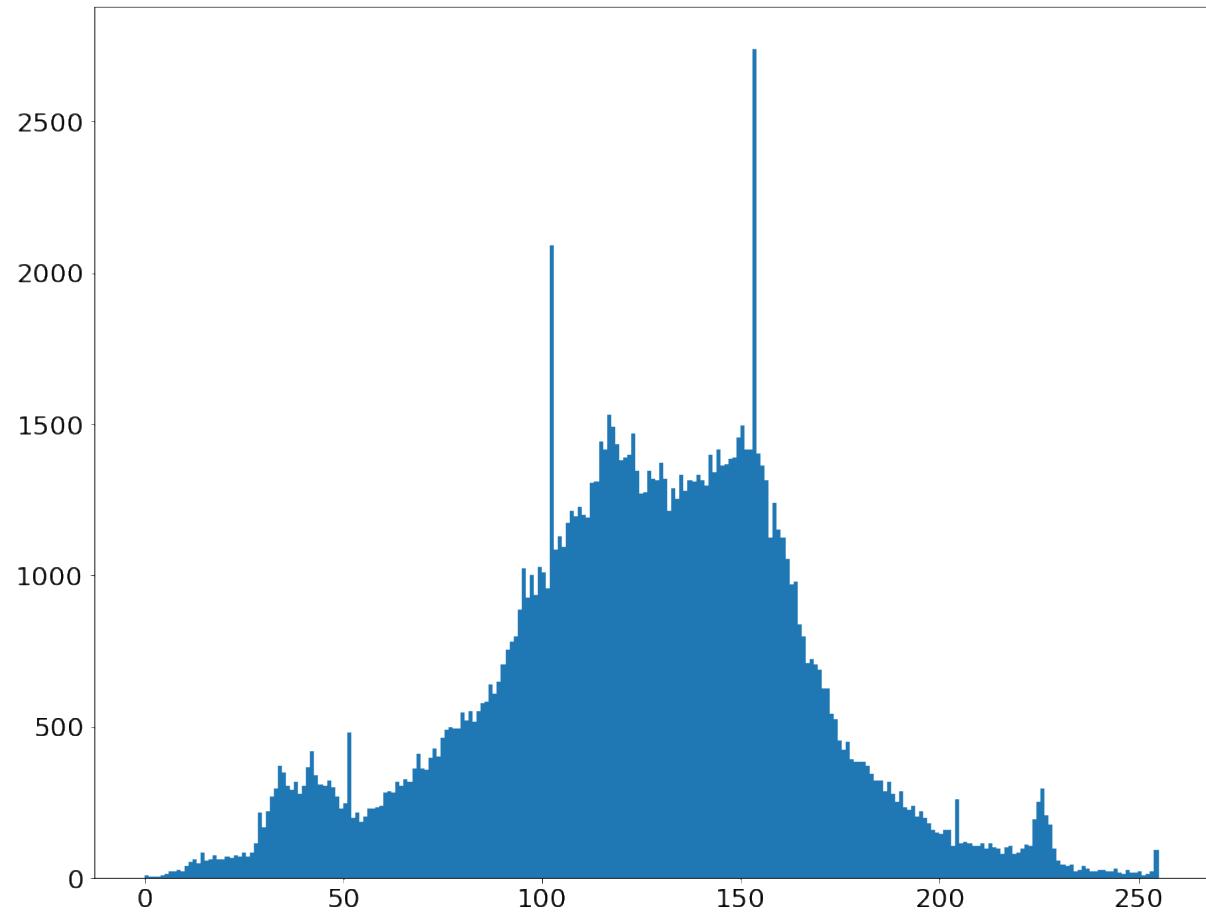
- Histogram of an image provides the frequency of the brightness (intensity) value in the image.

```
def histogram(im):
    h = np.zeros(255)
    for row in im.shape[0]:
        for col in im.shape[1]:
            val = im[row, col]
            h[val] += 1
```



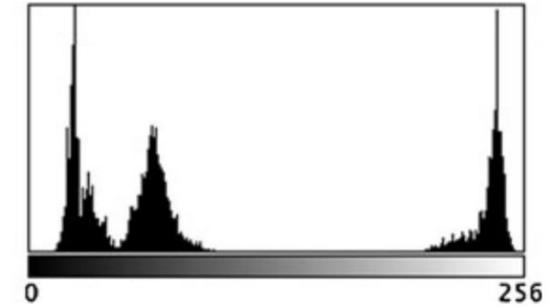
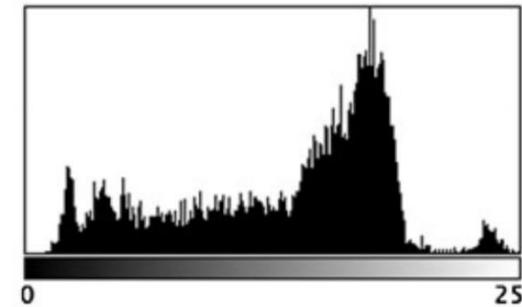
Histogram

- Histogram captures the distribution of gray levels in the image.
- How frequently each gray level occurs in the image





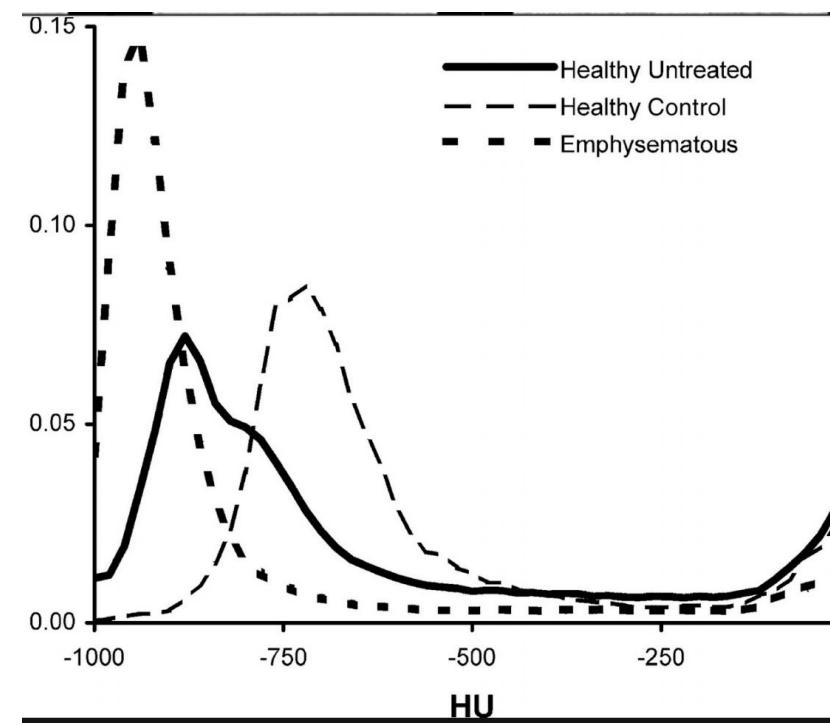
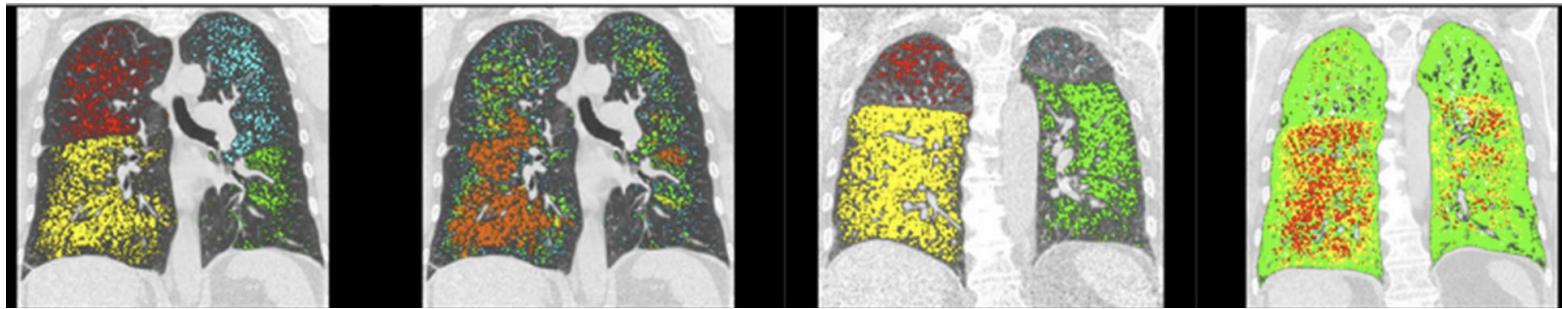
Histogram



Slide credit: Dr. Mubarak Shah



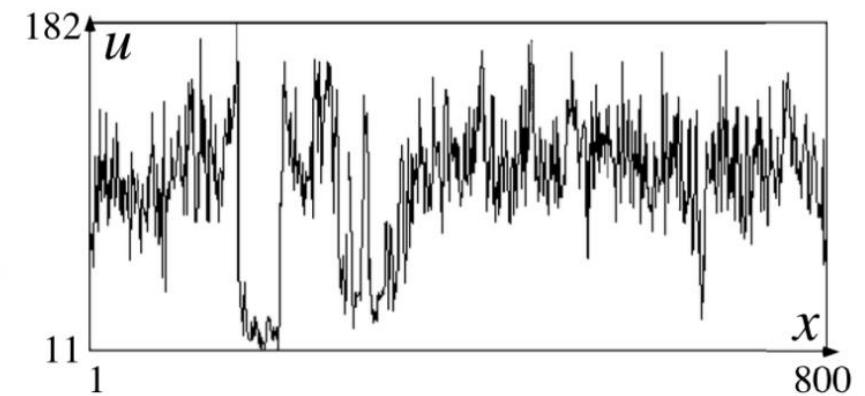
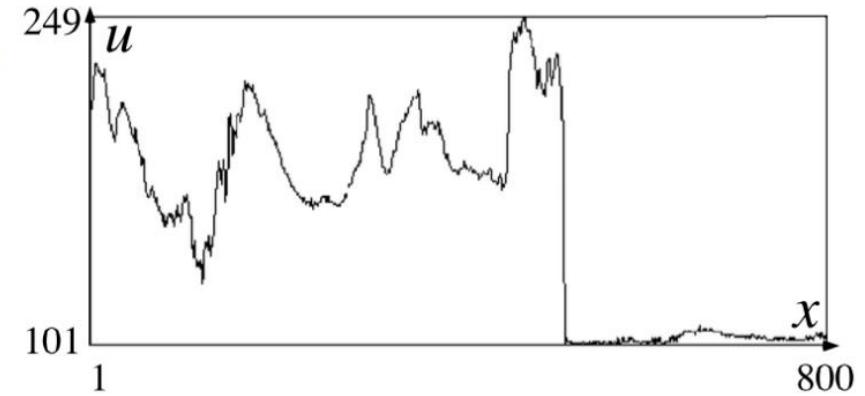
Histogram – use case



Slide credit: Dr. Mubarak Shah



Histogram – another use case





What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



Images as discrete functions

- Images are usually **digital (discrete)**:
 - Sample the 2D space on a regular grid
- Represented as a matrix of integer values

pixel

62	79	23	119	120	05	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30



Images as coordinates

Cartesian coordinates

$$f[n, m] = \begin{bmatrix} & \ddots & & \vdots \\ \cdots & f[-1, -1] & f[0, -1] & f[1, -1] \\ \cdots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] & \cdots \\ & f[-1, 1] & f[0, 1] & f[1, 1] & \ddots \end{bmatrix}$$

Notation for discrete functions



Images as coordinates

Cartesian coordinates

$$f[n, m] = \begin{bmatrix} & \ddots & & \vdots & \\ \dots & f[-1, -1] & f[0, -1] & f[1, -1] & \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] & \dots \\ f[-1, 1] & f[0, 1] & f[1, 1] & & \ddots \\ & \vdots & & & \end{bmatrix}$$

Notation for discrete functions





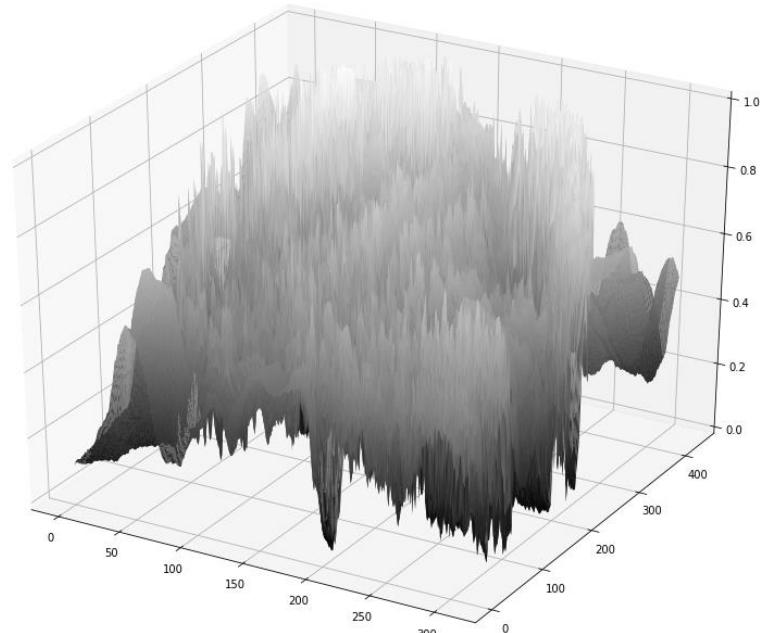
Images as functions

- An Image as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$f: [a,b] \times [c,d] \rightarrow [0,255]$$

Domain
support

range





Images as functions

- An Image as a function f from \mathbb{R}^2 to \mathbb{R}^M :

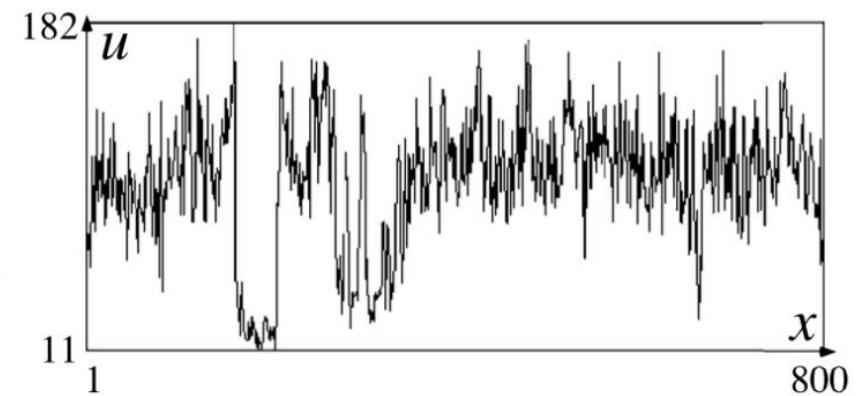
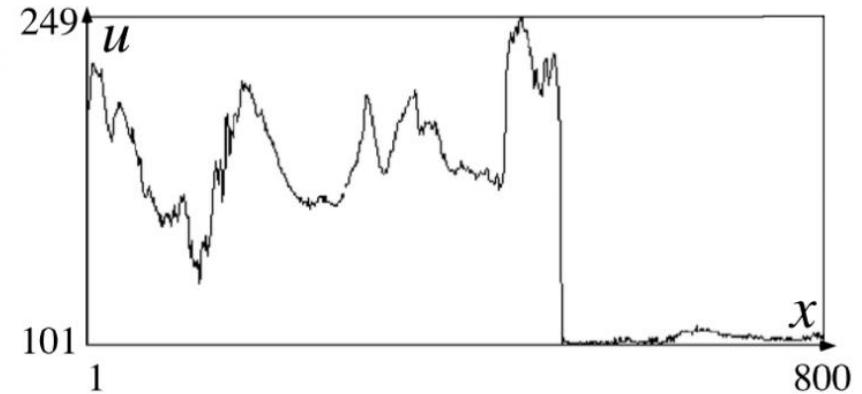
- $f(x, y)$ gives the **intensity** at position (x, y)
- Defined over a rectangle, with a finite range:

$$f: \underbrace{[a,b] \times [c,d]}_{\text{Domain support}} \rightarrow \underbrace{[0,255]}_{\text{range}}$$

- A color image: $f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$



Histograms are a type of image function





What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- **Linear systems (filters)**
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



Systems and Filters

Filtering:

- Forming a new image whose pixel values are transformed from original pixel values

Goals:

- Goal is to extract useful information from images, or transform images into another domain where we can modify/enhance image properties
 - Features (edges, corners, blobs...)
 - super-resolution; in-painting; de-noising



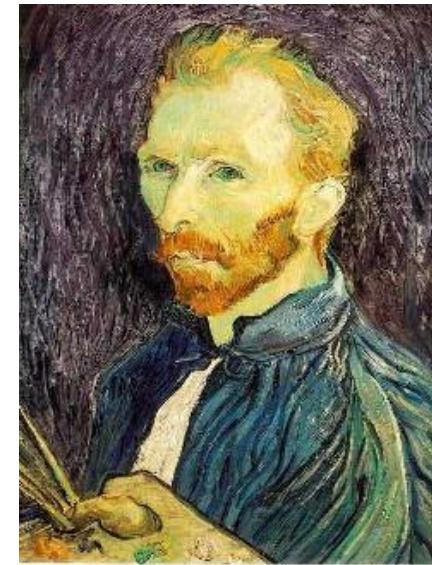
De-noising



Salt and pepper noise



Super-resolution



In-painting



Bertamio et al



System and Filters

- we define a system as a unit that converts an input function $f[n,m]$ into an output (or response) function $g[n,m]$, where (n,m) are the independent variables.
 - In the case for images, (n,m) represents the **spatial position in the image**.

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$



Images as coordinates

Cartesian coordinates

$$f[n, m] = \begin{bmatrix} \dots & f[-1, -1] & f[0, -1] & f[1, -1] \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] \\ f[-1, 1] & f[0, 1] & f[1, 1] & \dots \\ \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

Notation for discrete functions





2D discrete-space systems (filters)

S is the **system operator**, defined as a mapping or assignment of a member of the set of possible outputs $g[n,m]$ to each member of the set of possible inputs $f[n,m]$.

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$

$$g = \mathcal{S}[f], \quad g[n, m] = \mathcal{S}\{f[n, m]\}$$

$$f[n, m] \xrightarrow{S} g[n, m]$$



Filter example #1: Moving Average



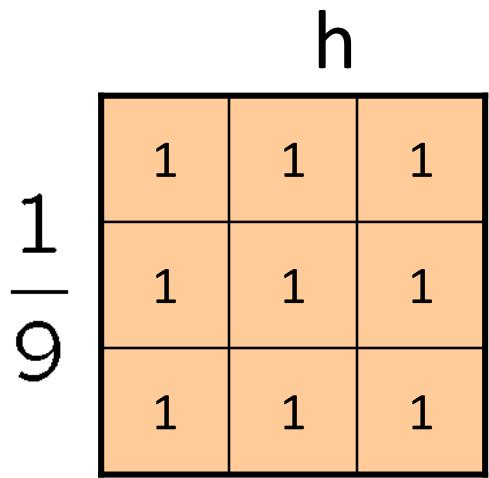


Filter example #1: Moving Average

2D DS moving average over a 3×3 window of neighborhood

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l]$$

$$= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$





Filter example #1: Moving Average

$f[n, m]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	90	90	90	90	90	0	0
0	0	0	0	90	90	90	90	90	0	0
0	0	0	0	90	90	90	90	90	0	0
0	0	0	0	90	0	90	90	90	0	0
0	0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[n, m]$

Courtesy of S. Seitz



Filter example #1: Moving Average

$$f[n, m]$$

$$g[n, m]$$



Filter example #1: Moving Average

$f[n, m]$

$$\mathbf{g}[n, m]$$



Filter example #1: Moving Average

$$f[n, m]$$

$$g[n, m]$$



Filter example #1: Moving Average

$$f[n, m]$$

$$g[n, m]$$



Filter example #1: Moving Average

$$f[n, m]$$

$$g[n, m]$$



Filter example #1: Moving Average

In summary:

- This filter creates a new pixel with an average of its neighborhood.
- Achieve smoothing effect (remove sharp features)

$$h[x, x]$$

1 | 9

1	1	1
1	1	1
1	1	1



Filter example #1: Moving Average



Filter example #2: Image Segmentation

- Image segmentation based on a simple threshold:

$$g[n, m] = \begin{cases} 255, & f[n, m] > 100 \\ 0, & \text{otherwise.} \end{cases}$$





Example Properties of systems

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$

- Stability

$$|f[n, m]| \leq k \implies |g[n, m]| \leq ck$$

- Invertibility

$$S^{-1}[S[f_i[n, m]]] = f[n, m]$$



Example Properties of linear systems

- Spatial properties

- Causality

for $n < n_0, m < m_0$, if $f[n, m] = 0 \implies g[n, m] = 0$

- Shift invariance:

$$f[n - n_0, m - m_0] \xrightarrow{\mathcal{S}} g[n - n_0, m - m_0]$$



These properties have to be true for **Linear Systems**

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$



These properties have to be true for **Linear SHIFT INVARIANT Systems**

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$

- Shift invariance:

$$f[n - n_0, m - m_0] \xrightarrow{\mathcal{S}} g[n - n_0, m - m_0]$$



What does shifting an image look like?

Cartesian coordinates

$$f[n, m] = \begin{bmatrix} \ddots & & \vdots & \\ & f[-1, 1] & f[0, 1] & f[1, 1] \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] \\ & f[-1, -1] & f[0, -1] & f[1, -1] \\ & & \vdots & \ddots \end{bmatrix}$$



Is the moving average system is shift invariant?

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l]$$

$f[n, m]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[n, m]$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
10	20	30	30	30	30	20	10			
10	10	10	0	0	0	0	0	0		



Is the moving average system is shift invariant?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$$

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

Let $f[n - n_0, m - m_0]$ be a shifted input of $f[n, m]$

Now let's pass $f[n - n_0, m - m_0]$ through the system:

$$\begin{aligned} f[n - n_0, m - m_0] &\xrightarrow{\mathcal{S}} \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - n_0 - k, m - m_0 - l] \\ &= g[n - n_0, m - m_0] \quad \text{Yes!} \end{aligned}$$



Is the moving average system causal?

$$g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$f[n, m]$ $g[n, m]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
10	20	30	30	30	30	30	20	10		
10	10	10	0	0	0	0	0	0		

for $n < n_0, m < m_0$, if $f[n, m] = 0 \implies g[n, m] = 0$



Linear Systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$

- Linear filtering:
 - Form a new image whose pixels are a weighted sum of original pixel values
 - Use the same set of weights at each point
- **S** is a linear system (function) iff it *S satisfies*

$$S[\alpha f_i[n, m] + \beta f_j[k, l]] = \alpha S[f_i[n, m]] + \beta S[f_j[k, l]]$$

superposition property



Linear Systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$

- Is the moving average a linear system?
- Is thresholding a linear system?
 - $f_1[n, m] + f_2[n, m] > T$
 - $f_1[n, m] < T$
 - $f_2[n, m] < T$

No!



How do we characterize a linear shift invariant (LSI) system?

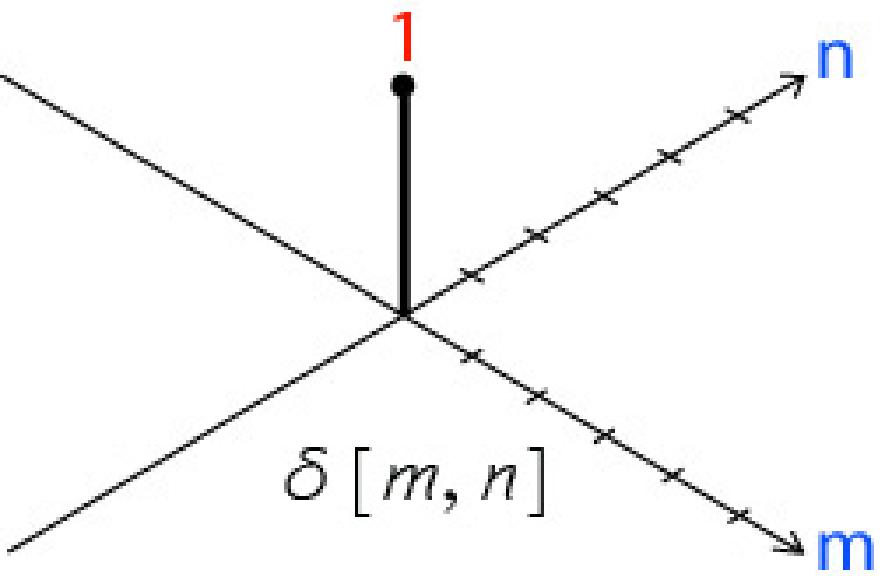
- Mathematically, how do you calculate $g[n,m]$ from $f[n,m]$

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$



2D impulse function

- 1 at [0,0].
- 0 everywhere else





Impulse response to the moving filter

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		?		
		$h[0,0]$		

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		1/9 h[0,0]	?	h[0,1]

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		1/9 h[0,0]	1/9 h[0,1]	
			?	h[1,1]

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		$1/9$ $h[0,0]$	$1/9$ $h[0,1]$	
			$1/9$ $h[1,1]$	

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		$1/9$ $h[0,0]$	$1/9$ $h[0,1]$? $h[0,2]$
			$1/9$ $h[1,1]$	

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

		$1/9$ $h[0,0]$	$1/9$ $h[0,1]$	0 $h[0,2]$
			$1/9$ $h[1,1]$	

$$\delta_2 \xrightarrow{s} h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response to the moving filter

0	0	0	0	0
0	$1/9$ $h[-1,-1]$	$1/9$	$1/9$	0
0	$1/9$	$1/9$ $h[0,0]$	$1/9$ $h[0,1]$	0 $h[0,2]$
0	$1/9$	$1/9$	$1/9$ $h[1,1]$	0
0	0	0	0	0

$$\delta_2 \xrightarrow{s} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$



Impulse response of the 3 by 3 moving average filter

$$h[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 \delta_2[n - k, m - l]$$

$$= \begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

$$\begin{matrix} h \\ \hline 1 & | & 0 \end{matrix} \quad \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Any linear shift invariant system

- By passing an impulse into a linear system, we get it's impulse response.
- So, if we don't know what the linear system is doing, we can pass an impulse into it to get a filter $h[n, m]$ that tells us what the system is actually doing.

$$\delta_2[n, m] \rightarrow \boxed{\text{System } S} \rightarrow h[n, m]$$

- But how do we use $h[n, m]$ to calculate $g[n, m]$ from $f[n, m]$

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$



General linear shift invariant system

- Let's say our input f is a 3x3 image:

$f[0,0]$	$f[0,1]$	$f[1,1]$
$f[1,0]$	$f[1,1]$	$f[1,2]$
$f[2,0]$	$f[2,1]$	$f[2,2]$

We can rewrite $f[n, m]$ as a sum of delta functions:

$$\begin{aligned}f[n, m] &= f[0,0] \times \delta_2[n - 0, m - 0] \\&+ f[0,1] \times \delta_2[n - 0, m - 1] \\&+ f[0,2] \times \delta_2[n - 0, m - 2] \\&+ \dots\end{aligned}$$

Or you can write it as: $f[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times \delta_2[n - k, m - l]$



3 properties we need:

Now, we know what happens when we send a delta function through a LSI system:

$$\delta_2[n, m] \rightarrow \boxed{\text{System } S} \rightarrow h[n, m]$$

We also know that LSI systems shift the output if the input is shifted:

$$\delta_2[n - k, m - l] \rightarrow \boxed{\text{System } S} \rightarrow h[n - k, m - l]$$

Finally, the superposition principle:

$$S[\alpha f_i[n, m] + \beta f_j[k, l]] = \alpha S[f_i[n, m]] + \beta S[f_j[k, l]]$$



We can generalize this superposition principle...

$$S[\alpha f_i[n, m] + \beta f_j[k, l]] = \alpha S[f_i[n, m]] + \beta S[f_j[k, l]]$$

... with our summation of deltas...

$$f[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times \delta_2[n - k, m - l]$$

... as follows:

$$\begin{aligned} & S \left[\sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times \delta_2[n - k, m - l] \right] \\ &= \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times S[\delta_2[n - k, m - l]] \end{aligned}$$



Linear shift invariant system

$$\begin{aligned} S \left[\sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times \delta_2[n - k, m - l] \right] \\ = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times S[\delta_2[n - k, m - l]] \\ = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] \times h[n - k, m - l] \end{aligned}$$



Linear Shift Invariant systems

An LSI system is completely specified by its impulse response.

$$\begin{aligned} f[n, m] &\rightarrow \boxed{\mathcal{S} \text{ LSI}} \rightarrow \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l] \\ \delta_2[n, m] &\rightarrow \boxed{\mathcal{S}} \rightarrow h[n, m] \end{aligned}$$

Discrete convolution

$$f[n, m] * h[n, m]$$



What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

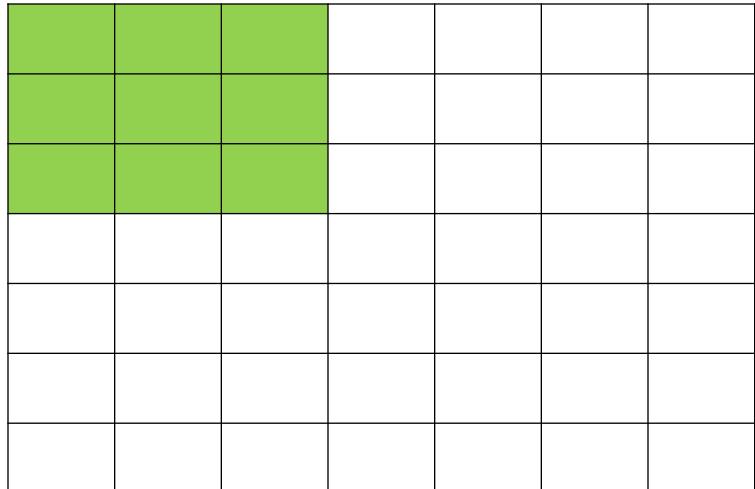


2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.



2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$

Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.



2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$

Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.



2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$

Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

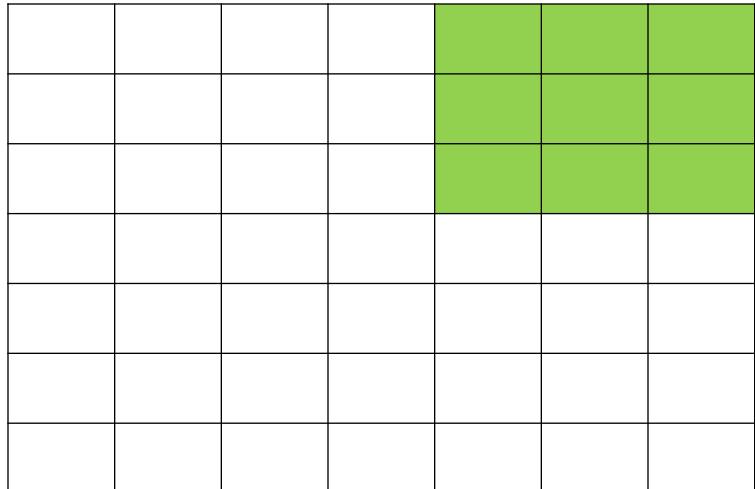


2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.



2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$

Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.



Convolution

$$f * h = \sum_k \sum_l f(k, l)h(-k, -l)$$

f = Image

h = Kernel

f

f_1	f_2	f_3
f_4	f_5	f_6
f_7	f_8	f_9

*

h_9	h_8	h_7
h_6	h_5	h_4
h_3	h_2	h_1

h_7	h_8	h_9
h_4	h_5	h_6
h_1	h_2	h_3

$X - flip$

h_1	h_2	h_3
h_4	h_5	h_6
h_7	h_8	h_9

$Y - flip$

$$\begin{aligned} f * h = & f_1 h_9 + f_2 h_8 + f_3 h_7 \\ & + f_4 h_6 + f_5 h_5 + f_6 h_4 \\ & + f_7 h_3 + f_8 h_2 + f_9 h_1 \end{aligned}$$



2D convolution example

1	2	3
4	5	6
7	8	9

Input

-1	0	1
-1	-2	-1
0	0	0

Kernel

-13	-20	-17
-18	-24	-18
13	20	17

Output



2D convolution example

1	2	1		
0	0	0	2	3
-1	-2	-1	4	5
7	8	9		

$$\begin{aligned} &= x[-1,-1] \cdot h[1,1] + x[0,-1] \cdot h[0,1] + x[1,-1] \cdot h[-1,1] \\ &\quad + x[-1,0] \cdot h[1,0] + x[0,0] \cdot h[0,0] + x[1,0] \cdot h[-1,0] \\ &\quad + x[-1,1] \cdot h[1,-1] + x[0,1] \cdot h[0,-1] + x[1,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 0 \cdot 0 + 1 \cdot 0 + 2 \cdot 0 + 0 \cdot (-1) + 4 \cdot (-2) + 5 \cdot (-1) = -13 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn



2D convolution example

1	2	1
0	0	0
1	2	3
-1	-2	-1
4	5	6
7	8	9

$$\begin{aligned} &= x[0,-1] \cdot h[1,1] + x[1,-1] \cdot h[0,1] + x[2,-1] \cdot h[-1,1] \\ &\quad + x[0,0] \cdot h[1,0] + x[1,0] \cdot h[0,0] + x[2,0] \cdot h[-1,0] \\ &\quad + x[0,1] \cdot h[1,-1] + x[1,1] \cdot h[0,-1] + x[2,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 0 + 2 \cdot 0 + 3 \cdot 0 + 4 \cdot (-1) + 5 \cdot (-2) + 6 \cdot (-1) = -20 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn



2D convolution example

	1	2	1
1	0	0	0
4	-1	5	-2
7	8	9	

$$\begin{aligned} &= x[1,-1] \cdot h[1,1] + x[2,-1] \cdot h[0,1] + x[3,-1] \cdot h[-1,1] \\ &\quad + x[1,0] \cdot h[1,0] + x[2,0] \cdot h[0,0] + x[3,0] \cdot h[-1,0] \\ &\quad + x[1,1] \cdot h[1,-1] + x[2,1] \cdot h[0,-1] + x[3,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 2 \cdot 0 + 3 \cdot 0 + 0 \cdot 0 + 5 \cdot (-1) + 6 \cdot (-2) + 0 \cdot (-1) = -17 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn



2D convolution example

1	2	1	2	3
0	0	0	5	6
-1	-2	-1	8	9

$$\begin{aligned} &= x[-1,0] \cdot h[1,1] + x[0,0] \cdot h[0,1] + x[1,0] \cdot h[-1,1] \\ &\quad + x[-1,1] \cdot h[1,0] + x[0,1] \cdot h[0,0] + x[1,1] \cdot h[-1,0] \\ &\quad + x[-1,2] \cdot h[1,-1] + x[0,2] \cdot h[0,-1] + x[1,2] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 1 \cdot 2 + 2 \cdot 1 + 0 \cdot 0 + 4 \cdot 0 + 5 \cdot 0 + 0 \cdot (-1) + 7 \cdot (-2) + 8 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output



2D convolution example

1	2	1
1	2	3
0	0	0
4	5	6
-1	-2	-1
7	8	9

$$\begin{aligned} &= x[0,0] \cdot h[1,1] + x[1,0] \cdot h[0,1] + x[2,0] \cdot h[-1,1] \\ &\quad + x[0,1] \cdot h[1,0] + x[1,1] \cdot h[0,0] + x[2,1] \cdot h[-1,0] \\ &\quad + x[0,2] \cdot h[1,-1] + x[1,2] \cdot h[0,-1] + x[2,2] \cdot h[-1,-1] \\ &= 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 1 + 4 \cdot 0 + 5 \cdot 0 + 6 \cdot 0 + 7 \cdot (-1) + 8 \cdot (-2) + 9 \cdot (-1) = -24 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn



2D convolution example

1	2	3	1
4	0	0	0
7	-1	-2	-1

$$\begin{aligned} &= x[1,0] \cdot h[1,1] + x[2,0] \cdot h[0,1] + x[3,0] \cdot h[-1,1] \\ &\quad + x[1,1] \cdot h[1,0] + x[2,1] \cdot h[0,0] + x[3,1] \cdot h[-1,0] \\ &\quad + x[1,2] \cdot h[1,-1] + x[2,2] \cdot h[0,-1] + x[3,2] \cdot h[-1,-1] \\ &= 2 \cdot 1 + 3 \cdot 2 + 0 \cdot 1 + 5 \cdot 0 + 6 \cdot 0 + 0 \cdot 0 + 8 \cdot (-1) + 9 \cdot (-2) + 0 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output



Convolution in 2D - examples



*

•0	•0	•0
•0	•1	•0
•0	•0	•0

=

?



Convolution in 2D - examples



Original

*

•0	•0	•0
•0	•1	•0
•0	•0	•0

=



Filtered
(no change)



Convolution in 2D - examples



*

•0	•0	•0
•0	•0	•1
•0	•0	•0

=

?



Convolution in 2D - examples



Original

*

•0	•0	•0
•0	•0	•1
•0	•0	•0

=



Shifted right
By 1 pixel



Convolution in 2D - examples



Original

$$\begin{matrix} \text{Original} \\ \ast \\ \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} \end{matrix}$$

=

?



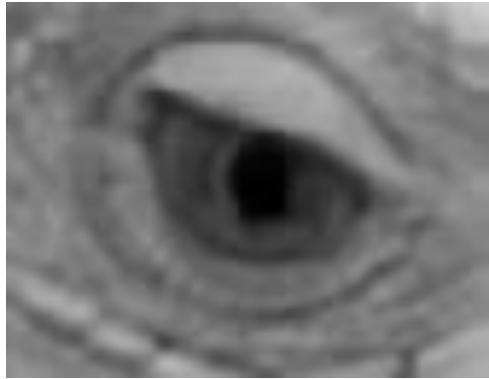
Convolution in 2D - examples



Original

$$\ast \frac{1}{9} \begin{array}{|ccc|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array}$$

=



Blur (with a
box filter)



Convolution in 2D - examples



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

–

1	–	9
•1	•1	•1
•1	•1	•1

= ?

(Note that filter sums to 1)

“details of the image”

•0	•0	•0
•0	•1	•0
•0	•0	•0

+

•0	•0	•0
•0	•1	•0
•0	•0	•0

–

1
–
9

•1	•1	•1
•1	•1	•1
•1	•1	•1



What does blurring take away?



-



II



- Let's add it back:



+



II





Convolution in 2D – Sharpening filter



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0



•1	•1	•1
•1	•1	•1
•1	•1	•1

$\frac{1}{9}$



Sharpening filter: Accentuates differences with local average



Image support and edge effect

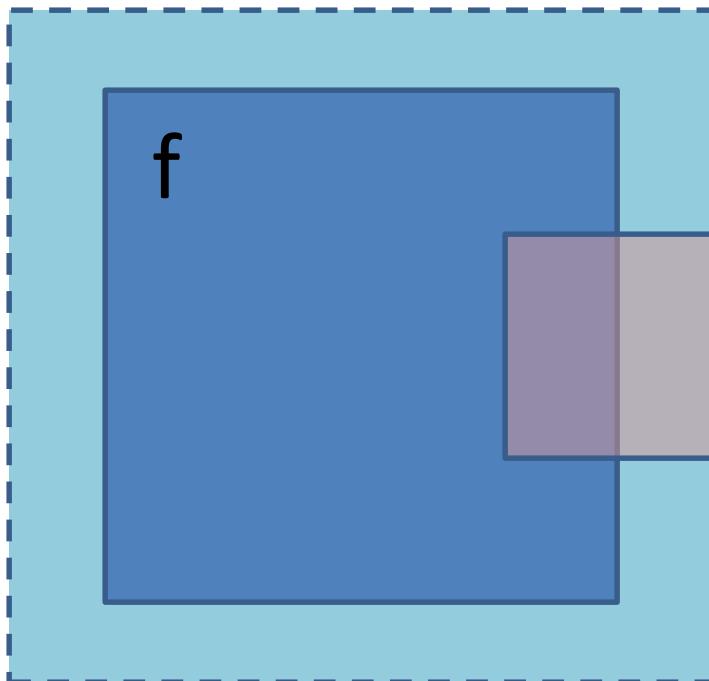
- A computer will only convolve **finite support signals**.
 - That is: images that are zero for n,m outside some rectangular region
 - numpy's convolution performs 2D DS convolution of finite-support signals.

$$\begin{array}{c} \text{Blue square: } N_1 \times M_1 \\ * \quad \text{Red square: } N_2 \times M_2 \\ = \quad \text{Result: } (N_1 + N_2 - 1) \times (M_1 + M_2 - 1) \end{array}$$
A diagram illustrating a convolution operation. On the left, there is a large blue square labeled $N_1 \times M_1$. To its right is a smaller red square labeled $N_2 \times M_2$, with an asterisk (*) indicating multiplication. Between the two squares is an equals sign (=). To the right of the equals sign is a larger green square divided into four quadrants. The top-left quadrant is red, matching the size of the red square. The other three quadrants are green and have a darker green border. This visualizes how a smaller filter kernel (red) is applied to a larger input image (blue) to produce a larger output feature map (green), with the output size being $(N_1 + N_2 - 1) \times (M_1 + M_2 - 1)$.

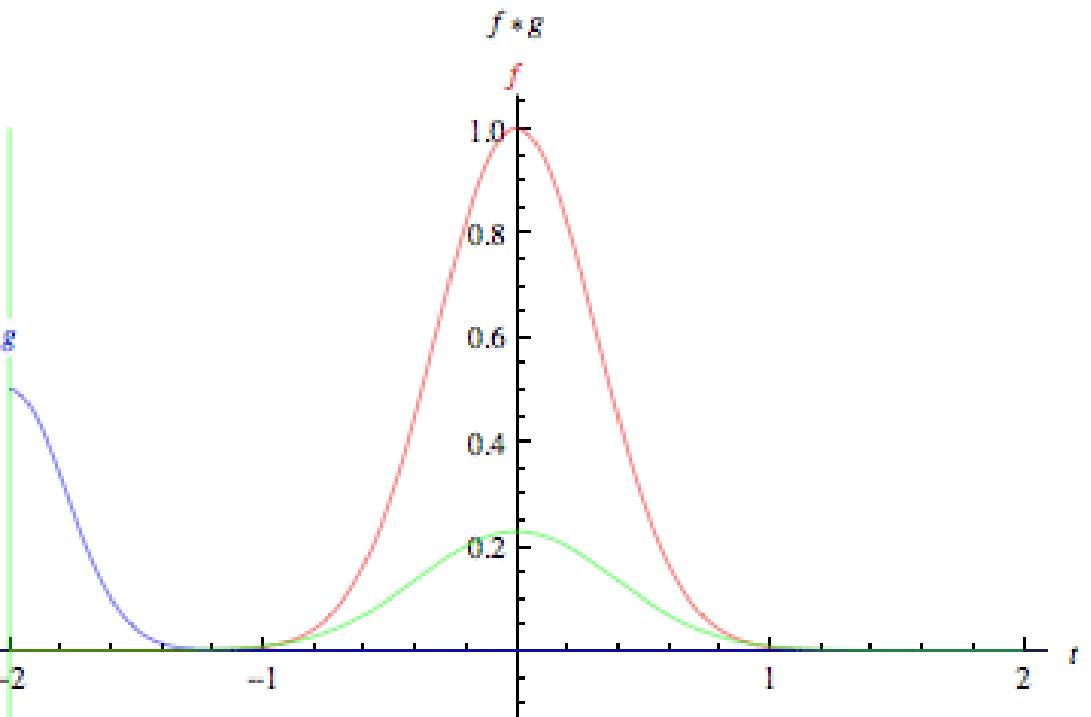
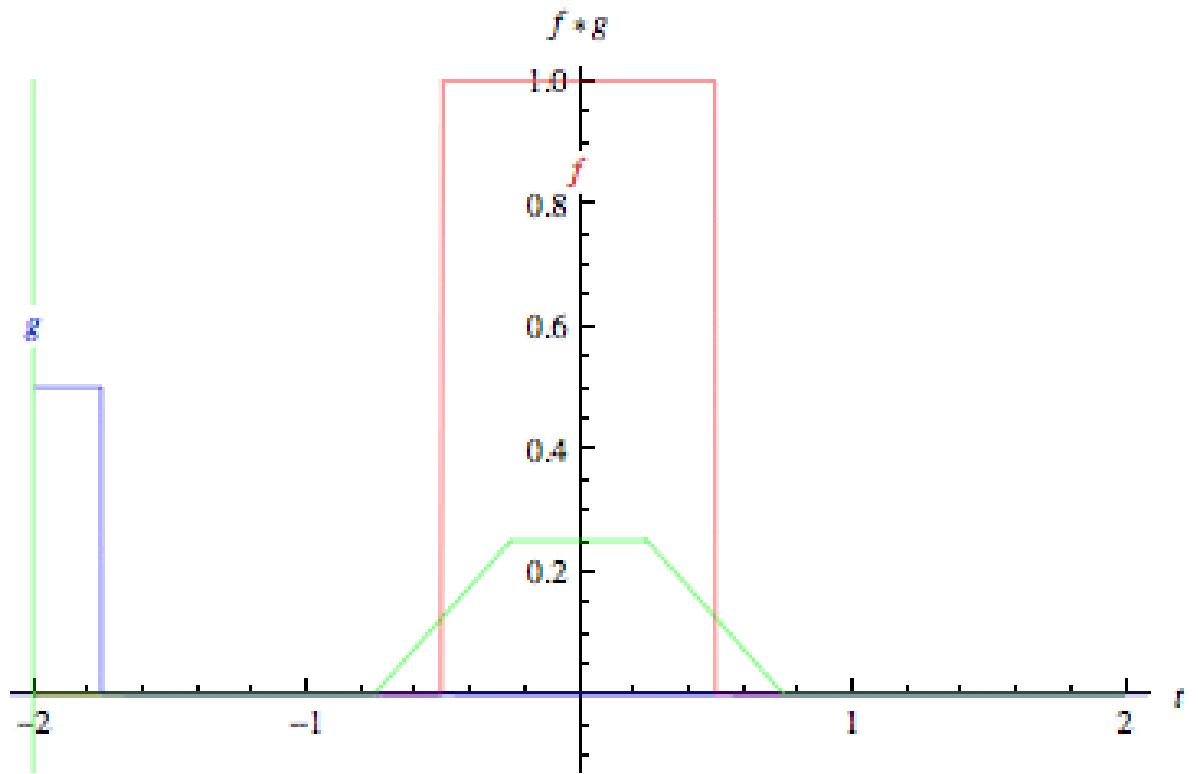


Image support and edge effect

- A computer will only convolve **finite support signals**.
- What happens at the edge?



- zero “padding”
- edge value replication
- mirror extension
- more (beyond the scope of this class)





What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7



(Cross) correlation (symbol: $\ast\ast$)

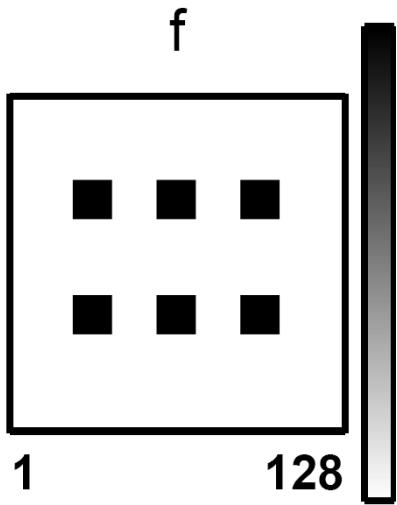
Cross correlation of two 2D signals $f[n,m]$ and $h[n,m]$:

$$f[n, m] \ast\ast h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[n, m] \times h[n + k, m + l]$$

Equivalent to a convolution without the flip

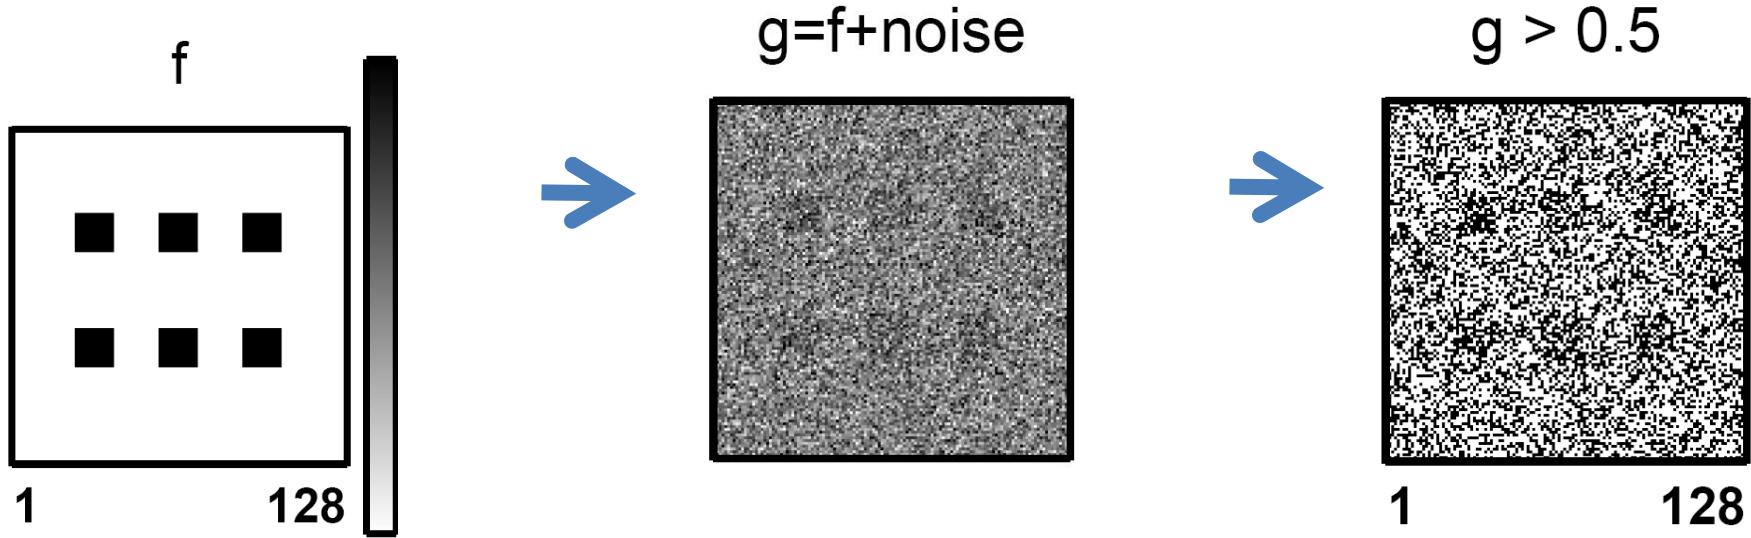


(Cross) correlation – example



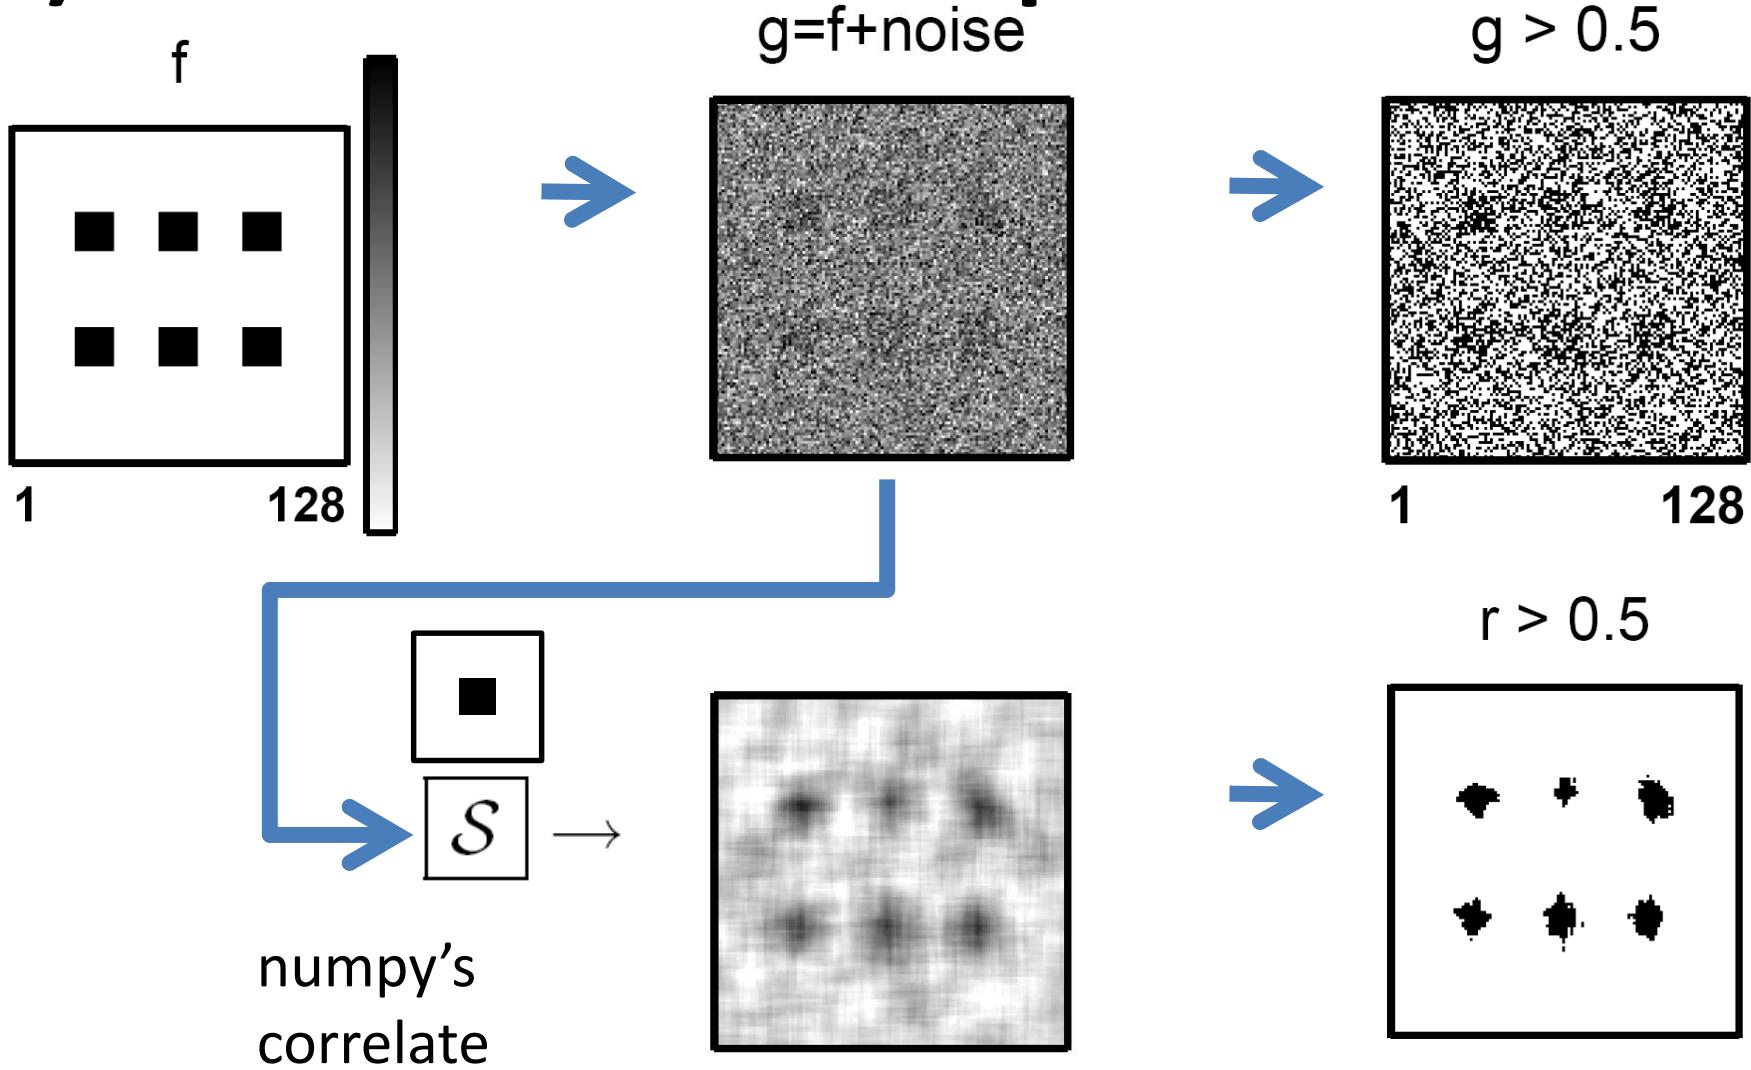


(Cross) correlation – example





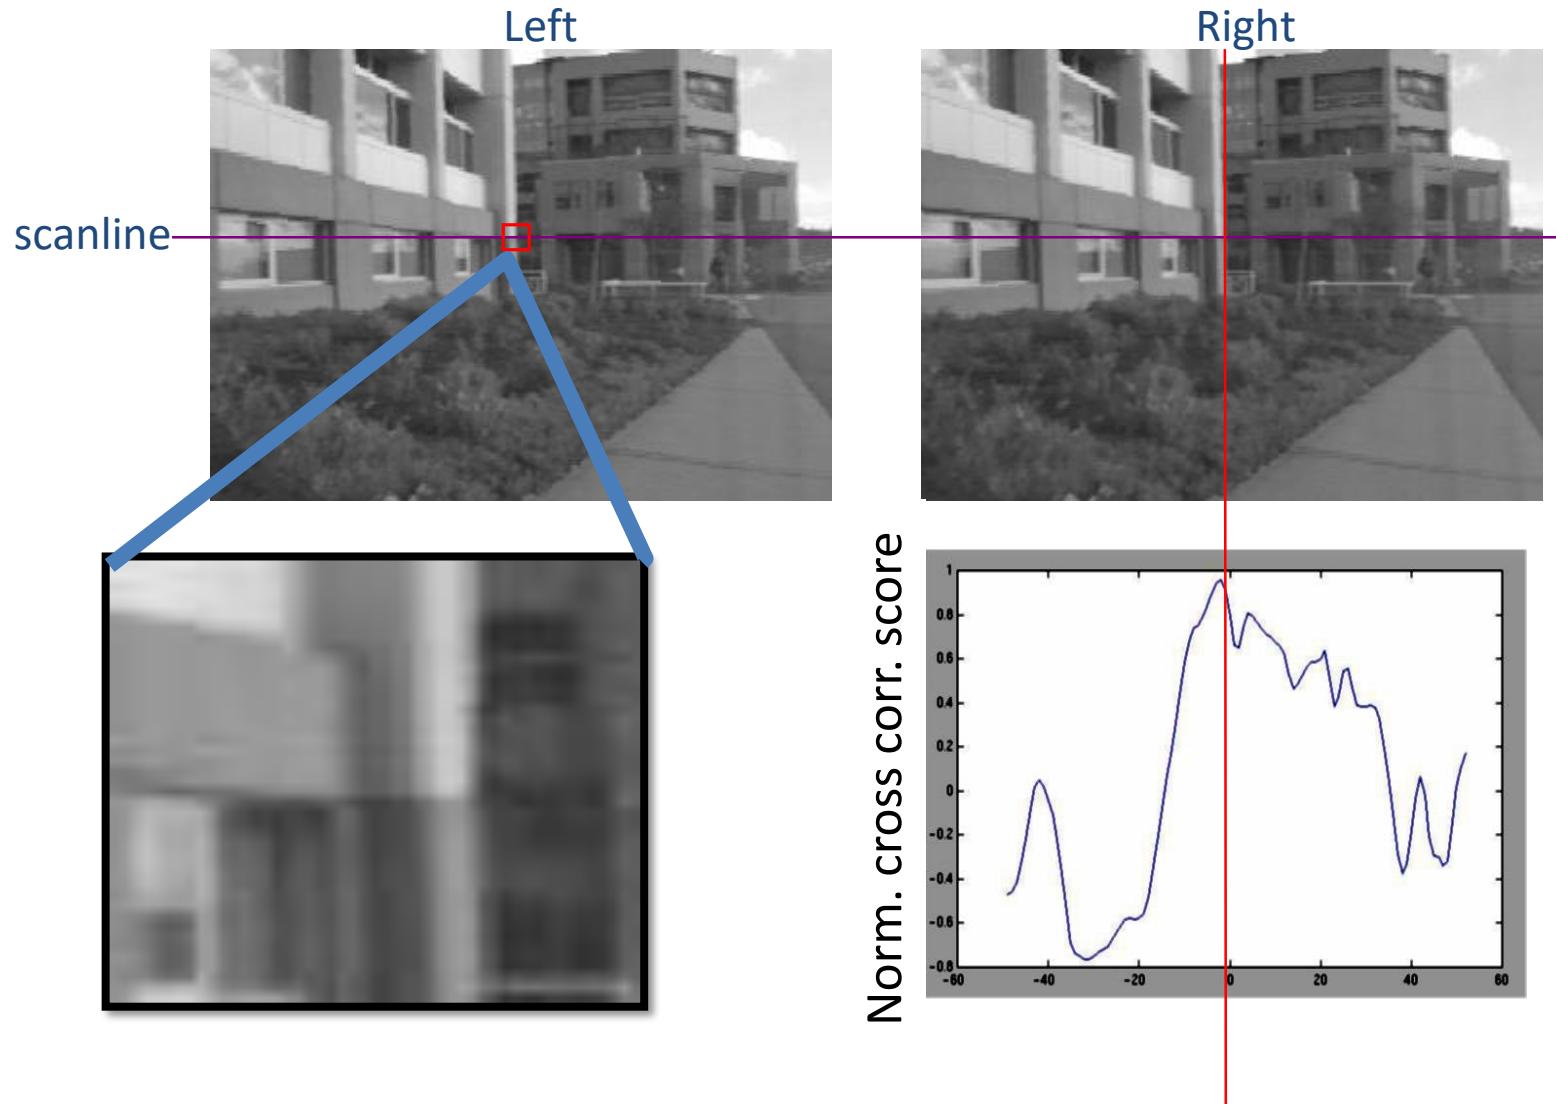
(Cross) correlation – example



Courtesy of J. Fessler

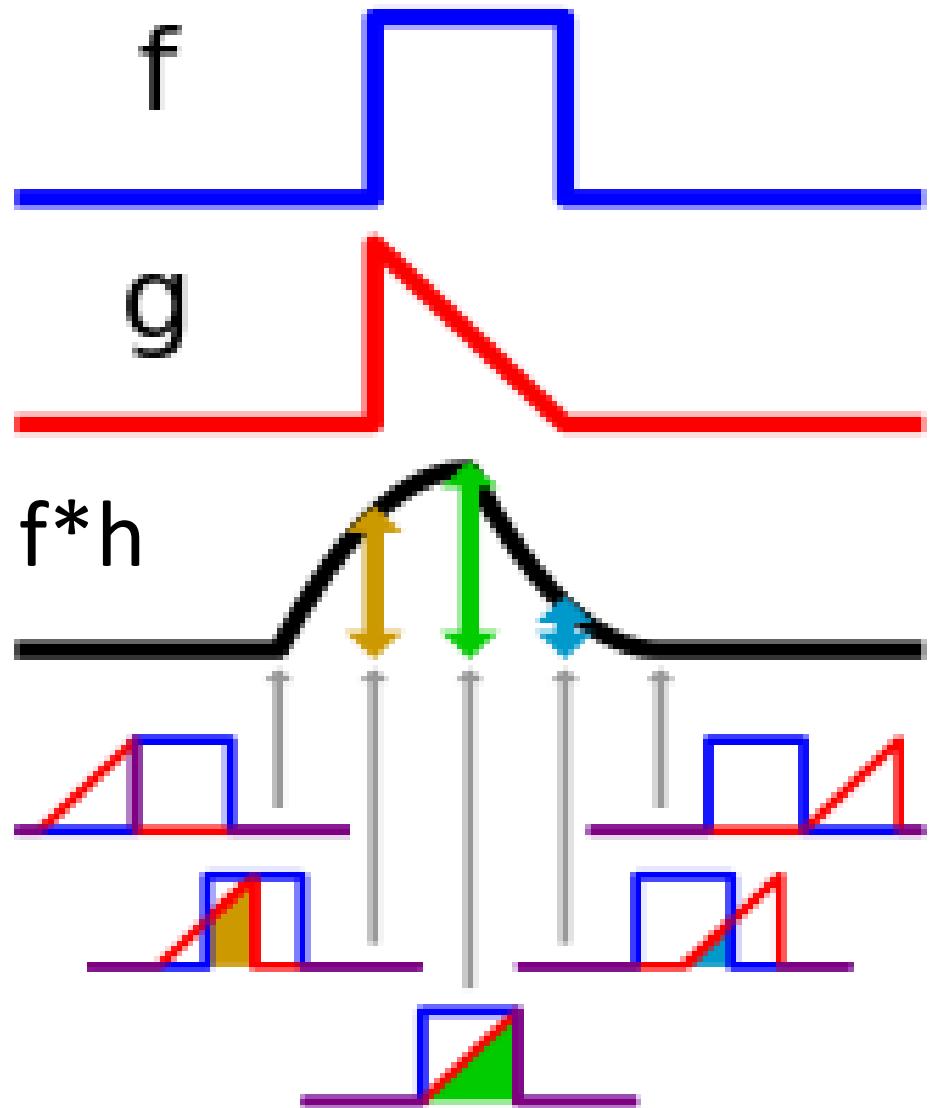


(Cross) correlation – example

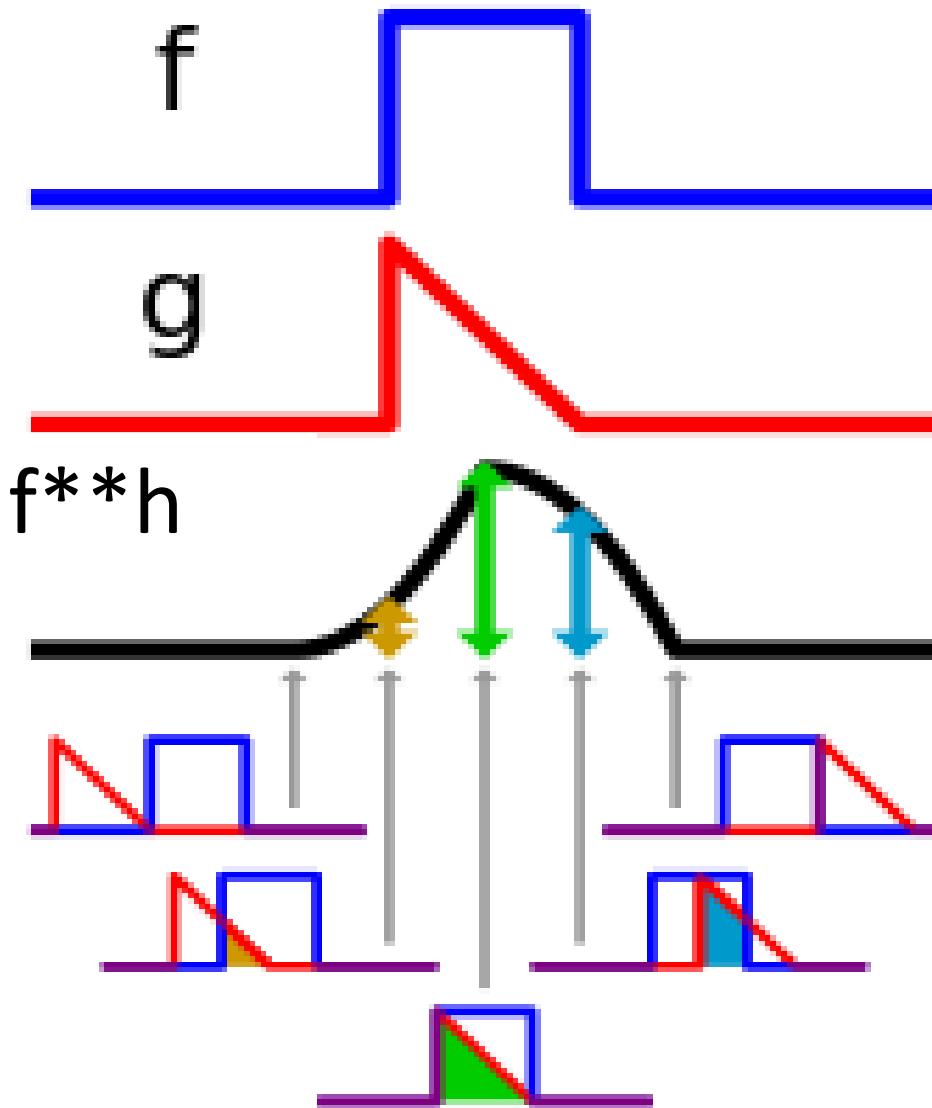




Convolution



Cross-correlation





Cross Correlation Application: Vision system for TV remote control

- uses template matching

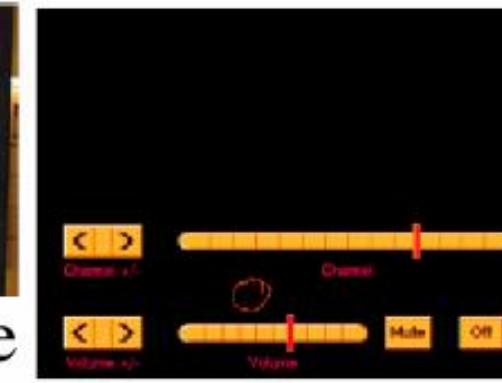
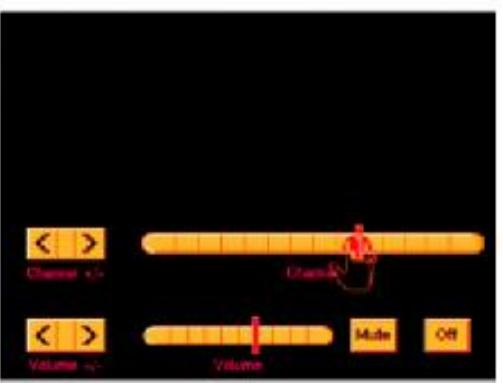
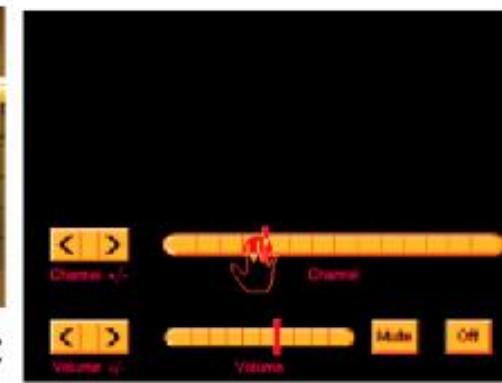
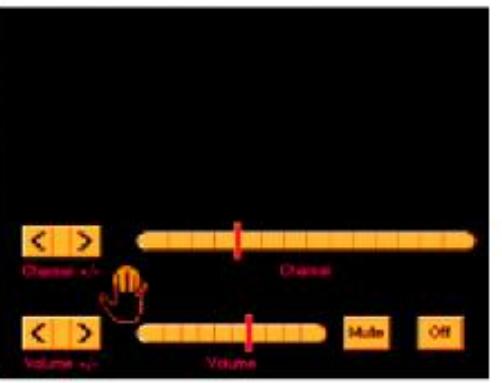


Figure from “Computer Vision for Interactive Computer Graphics,” W.Freeman et al, IEEE Computer Graphics and Applications, 1998 copyright 1998, IEEE



Properties

- Commutative property:

$$f \ast\ast h = h \ast\ast f$$

- Associative property:

$$(f \ast\ast h_1) \ast\ast h_2 = f \ast\ast (h_1 \ast\ast h_2)$$

- Distributive property:

$$f \ast\ast (h_1 + h_2) = (f \ast\ast h_1) + (f \ast\ast h_2)$$

The order doesn't matter! $h_1 \ast\ast h_2 = h_2 \ast\ast h_1$



Properties

- Shift property:

$$f[n, m] \ast\ast \delta_2[n - n_0, m - m_0] = f[n - n_0, m - m_0]$$

- Shift-invariance:

$$\begin{aligned} g[n, m] &= f[n, m] \ast\ast h[n, m] \\ &\implies f[n - l_1, m - l_1] \ast\ast h[n - l_2, m - l_2] \\ &= g[n - l_1 - l_2, m - l_1 - l_2] \end{aligned}$$



Convolution vs. (Cross) Correlation

- A **convolution** is an integral that expresses the amount of overlap of one function as it is shifted over another function.
 - convolution is a filtering operation
- **Correlation** compares the *similarity of two sets of data*. Correlation computes a measure of similarity of two input signals as they are shifted by one another. The correlation result reaches a maximum at the time when the two signals match best .
 - correlation is a measure of relatedness of two signals



What we have learned today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation