

Optical Flow and Feature Tracking

Many slides adapted from S. Seitz, R. Szeliski, M. Pollefeys, S. Lazebnik

Video

- A video is a sequence of frames captured over time
- Now our image data is a function of space (x, y) and time (t)




Motion and perceptual organization

- Sometimes, motion is foremost cue



Motion and perceptual organization


- Even “impoverished” motion data can evoke a strong percept



G. Johansson, “Visual Perception of Biological Motion and a Model For Its Analysis”, *Perception and Psychophysics* 14, 201-211, 1973.

Motion and perceptual organization

- Even “impoverished” motion data can evoke a strong percept





G. Johansson, “Visual Perception of Biological Motion and a Model For Its Analysis”, *Perception and Psychophysics* 14, 201-211, 1973.

Uses of motion

- Estimating 3D structure
- Segmenting objects based on motion cues
- Learning dynamical models
- Recognizing events and activities
- Improving video quality (motion stabilization)


Motion field

- The motion field is the projection of the 3D scene motion into the image



Motion field

- $\mathbf{P}(t)$ is a moving 3D point
- Velocity of scene point: $\mathbf{V} = d\mathbf{P}/dt$
- $\mathbf{p}(t) = (x(t), y(t))$ is the projection of \mathbf{P} in the image
- Apparent velocity \mathbf{v} in the image: given by components $v_x = dx/dt$ and $v_y = dy/dt$
- These components are known as the *motion field* of the image




Motion estimation techniques

- Direct methods
 - Directly recover image motion at each pixel from spatio-temporal image brightness variations
 - Dense motion fields, but sensitive to appearance variations
 - Suitable for video and when image motion is small
- Feature-based methods
 - Extract visual features (corners, textured areas) and track them over multiple frames
 - Sparse motion fields, but more robust tracking
 - Suitable when image motion is large

Optical flow



- Definition: optical flow is the *apparent* motion of brightness patterns in the image
- Ideally, optical flow would be the same as the motion field
- Have to be careful: apparent motion can be caused by lighting changes without any actual motion
 - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination

Optical Flow \neq Motion Field



Motion field exists but no optical flow

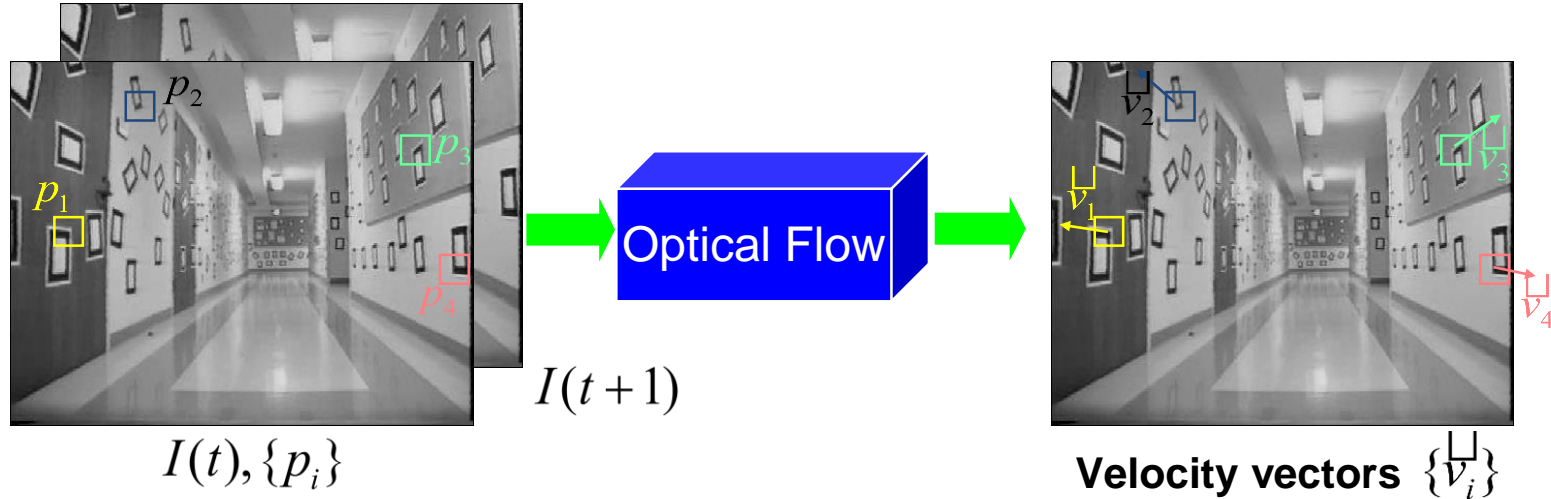
(a)



No motion field but shading changes

(b)

What is Optical Flow?



Optical flow is the relation of the motion field


- *the 2D projection of the physical movement of points relative to the observer to 2D displacement of pixel patches on the image plane.*

Common assumption:


The appearance of the image patches do not change (brightness constancy)

$$I(p_i, t) = I(p_i + \underline{v}_i, t + 1)$$

Estimating optical flow




$I(x,y,t-1)$



$I(x,y,t)$

- Given two subsequent frames, estimate the apparent motion field between them.
- Key assumptions
 - **Brightness constancy:** projection of the same point looks the same in every frame
 - **Small motion:** points do not move very far
 - **Spatial coherence:** points move like their neighbors

Optical Flow Assumptions: Brightness Constancy




Assumption

Image measurements (e.g. brightness) in a small region remain the same although their location may change.

$$I(x + u, y + v, t + 1) = I(x, y, t)$$

(assumption)


Optical Flow Assumptions: Spatial Coherence



Assumption

- * Neighboring points in the scene typically belong to the same surface and hence typically have similar motions.
- * Since they also project to nearby points in the image, we expect spatial coherence in image flow.

Optical Flow Assumptions: Temporal Persistence



Assumption:

The image motion of a surface patch changes gradually over time.


Estimating Optical Flow

- Assume the image intensity I is constant

Time = t



Time = $t + dt$



$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

Brightness Constancy Equation

$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

First order Taylor Expansion

$$= I(x, y, t) + \frac{\partial I}{\partial x} dx + \frac{\partial I}{\partial y} dy + \frac{\partial I}{\partial t} dt$$

Simplify notations:

$$I_x dx + I_y dy + I_t dt = 0$$


Divide by dt and denote:

$$u = \frac{dx}{dt} \quad v = \frac{dy}{dt}$$

$$I_x u + I_y v = -I_t$$


Problem I: One equation, two unknowns

The aperture problem




Perceived motion

The aperture problem




The barber pole illusion



http://en.wikipedia.org/wiki/Barberpole_illusion

The barber pole illusion



http://en.wikipedia.org/wiki/Barberpole_illusion

The barber pole illusion



http://en.wikipedia.org/wiki/Barberpole_illusion

Solving the aperture problem (grayscale image)

- How to get more equations for a pixel?
- **Spatial coherence constraint:** pretend the pixel's neighbors have the same (u,v)
 - If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

$$\begin{matrix} A & d & = & b \\ 25 \times 2 & 2 \times 1 & & 25 \times 1 \end{matrix}$$

Solving the aperture problem

Prob: we have more equations than unknowns

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 \quad 25 \times 1 \end{matrix} \longrightarrow \text{minimize } \|Ad - b\|^2$$

Solution: solve least squares problem

- minimum least squares solution given by solution (in d) of:

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 \quad 2 \times 1 \end{matrix}$$

$$\begin{matrix} \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} = - & \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix} \\ A^T A & & A^T b \end{matrix}$$

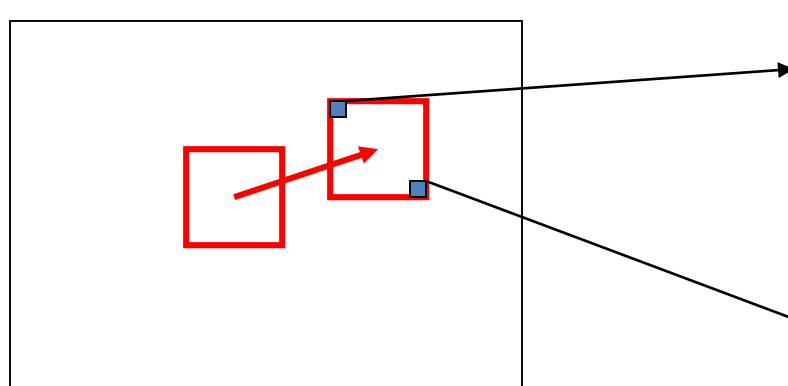
- The summations are over all pixels in the K x K window
- This technique was first proposed by Lucas & Kanade (1981)

Local smoothness

Lucas Kanade (1981)

$$I_x u + I_y v = -I_t \quad \Rightarrow \quad \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = -I_t$$

Assume constant (u,v) in small neighborhood


$$\begin{bmatrix} I_{x1} & I_{y1} \\ I_{x2} & I_{y2} \\ \vdots & \vdots \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_{t1} \\ I_{t2} \\ \vdots \end{bmatrix}$$

$$A \mathbf{u} = b$$

Lucas Kanade (1981)

Goal: Minimize $\|A\mathbf{u} - b\|^2$

Method: Least-Squares

$$A\mathbf{u} = b$$



$$\underbrace{A^T}_{2 \times 2} \underbrace{A}_{2 \times 1} \underbrace{\mathbf{u}}_{2 \times 1} = \underbrace{A^T b}_{2 \times 1}$$



$$\mathbf{u} = (A^T A)^{-1} A^T b$$

Conditions for solvability

Optimal (u, v) satisfies Lucas-Kanade equation

$$\underbrace{\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix}}_{A^T A} \begin{bmatrix} u \\ v \end{bmatrix} = - \underbrace{\begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}}_{A^T b}$$


When is this solvable? I.e., what are good points to track?

- $A^T A$ should be invertible
- $A^T A$ should not be too small due to noise
 - eigenvalues λ_1 and λ_2 of $A^T A$ should not be too small
- $A^T A$ should be well-conditioned
 - λ_1 / λ_2 should not be too large ($\lambda_1 =$ larger eigenvalue)

Does this remind you of anything?


Criteria for Harris corner detector

Edge





$$\sum \nabla I (\nabla I)^T$$

- large gradients, all the same
- large λ_1 , small λ_2




Low texture region



$$\sum \nabla I (\nabla I)^T$$


- gradients have small magnitude
- small λ_1 , small λ_2

High textured region



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large λ_1 , large λ_2



Optical flow




Errors in Lucas-Kanade

- The motion is large (larger than a pixel)
 - Iterative refinement
 - Coarse-to-fine estimation
 - Exhaustive neighborhood search (feature matching)
- A point does not move like its neighbors
 - Motion segmentation
- Brightness constancy does not hold
 - Exhaustive neighborhood search with normalized correlation

Iterative Refinement


- Estimate velocity at each pixel using one iteration of Lucas and Kanade estimation
- Warp one image toward the other using the estimated flow field
(easier said than done)
- Refine estimate by repeating the process

Optical Flow: Iterative Estimation




(using d for *displacement* here instead of u)


Optical Flow: Iterative Estimation




Optical Flow: Iterative Estimation



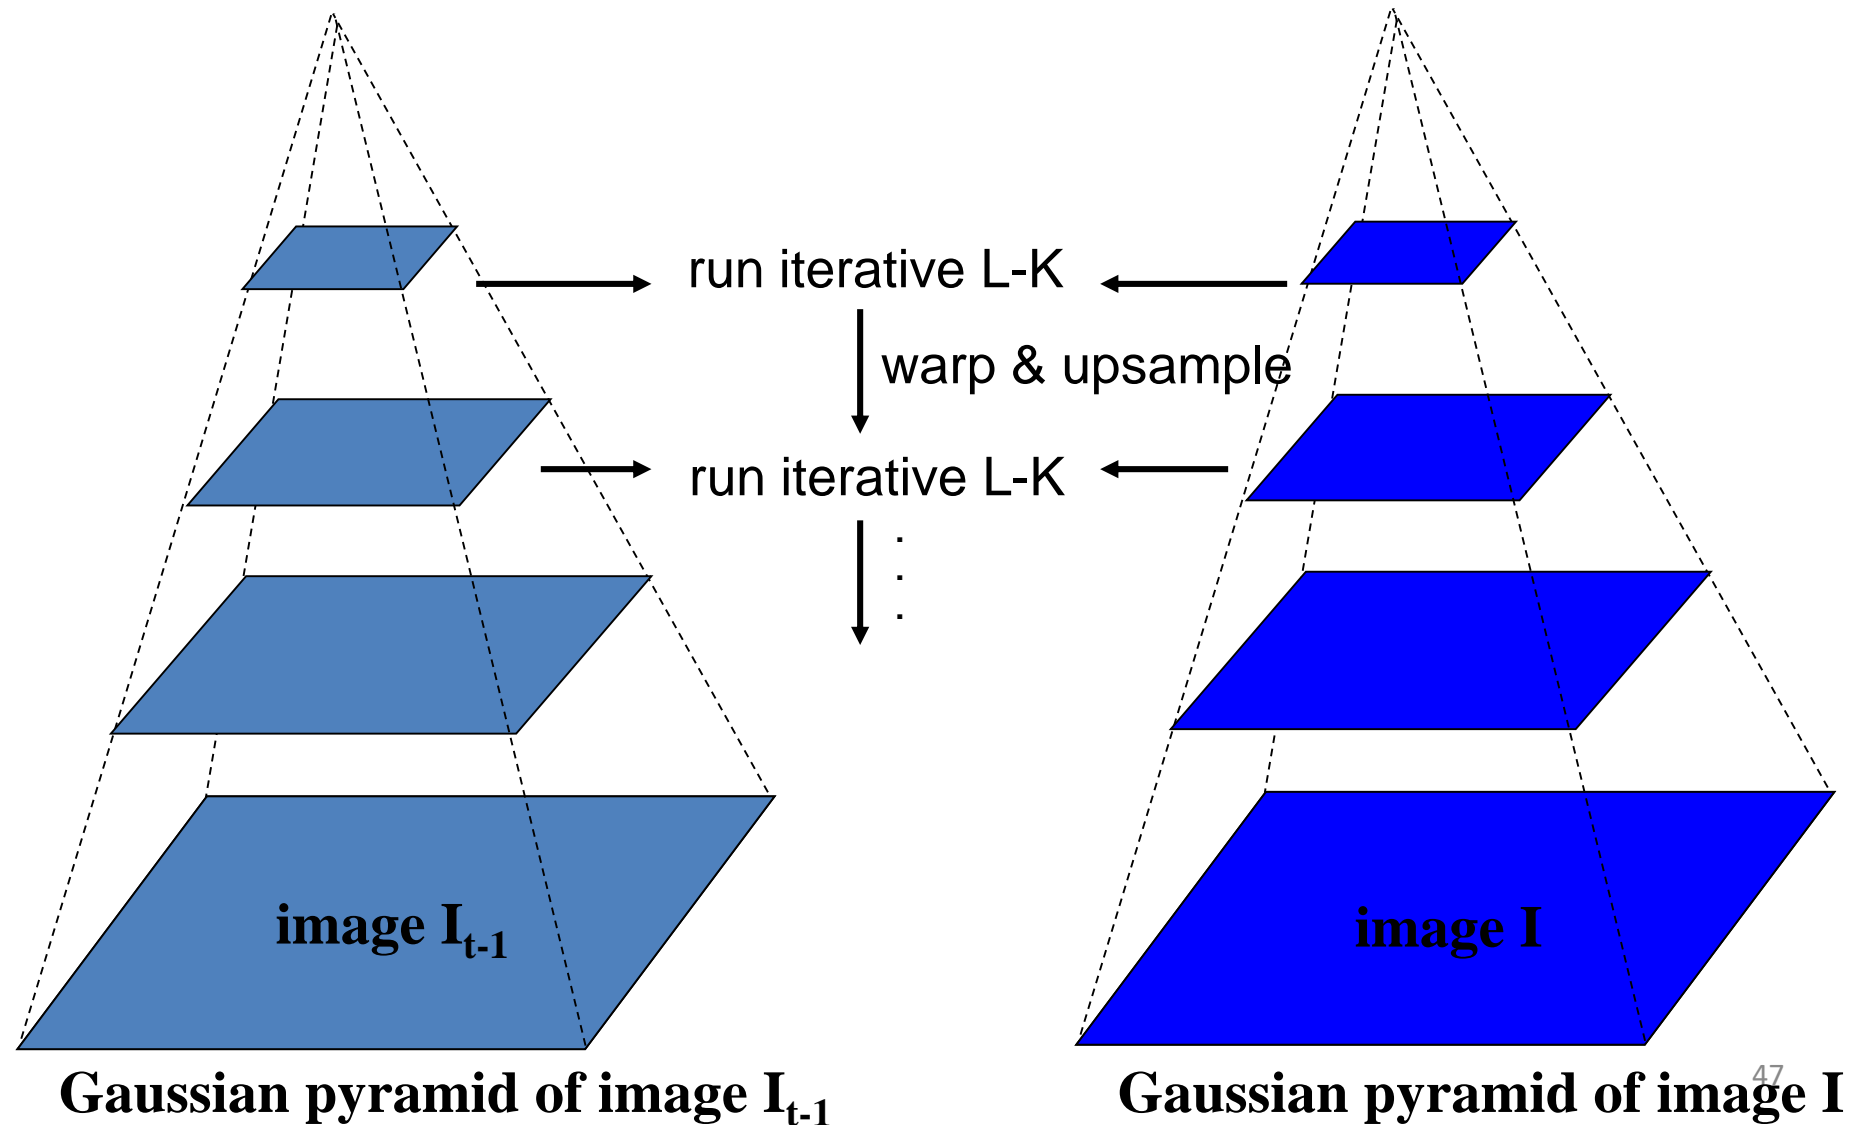
Optical Flow: Iterative Estimation



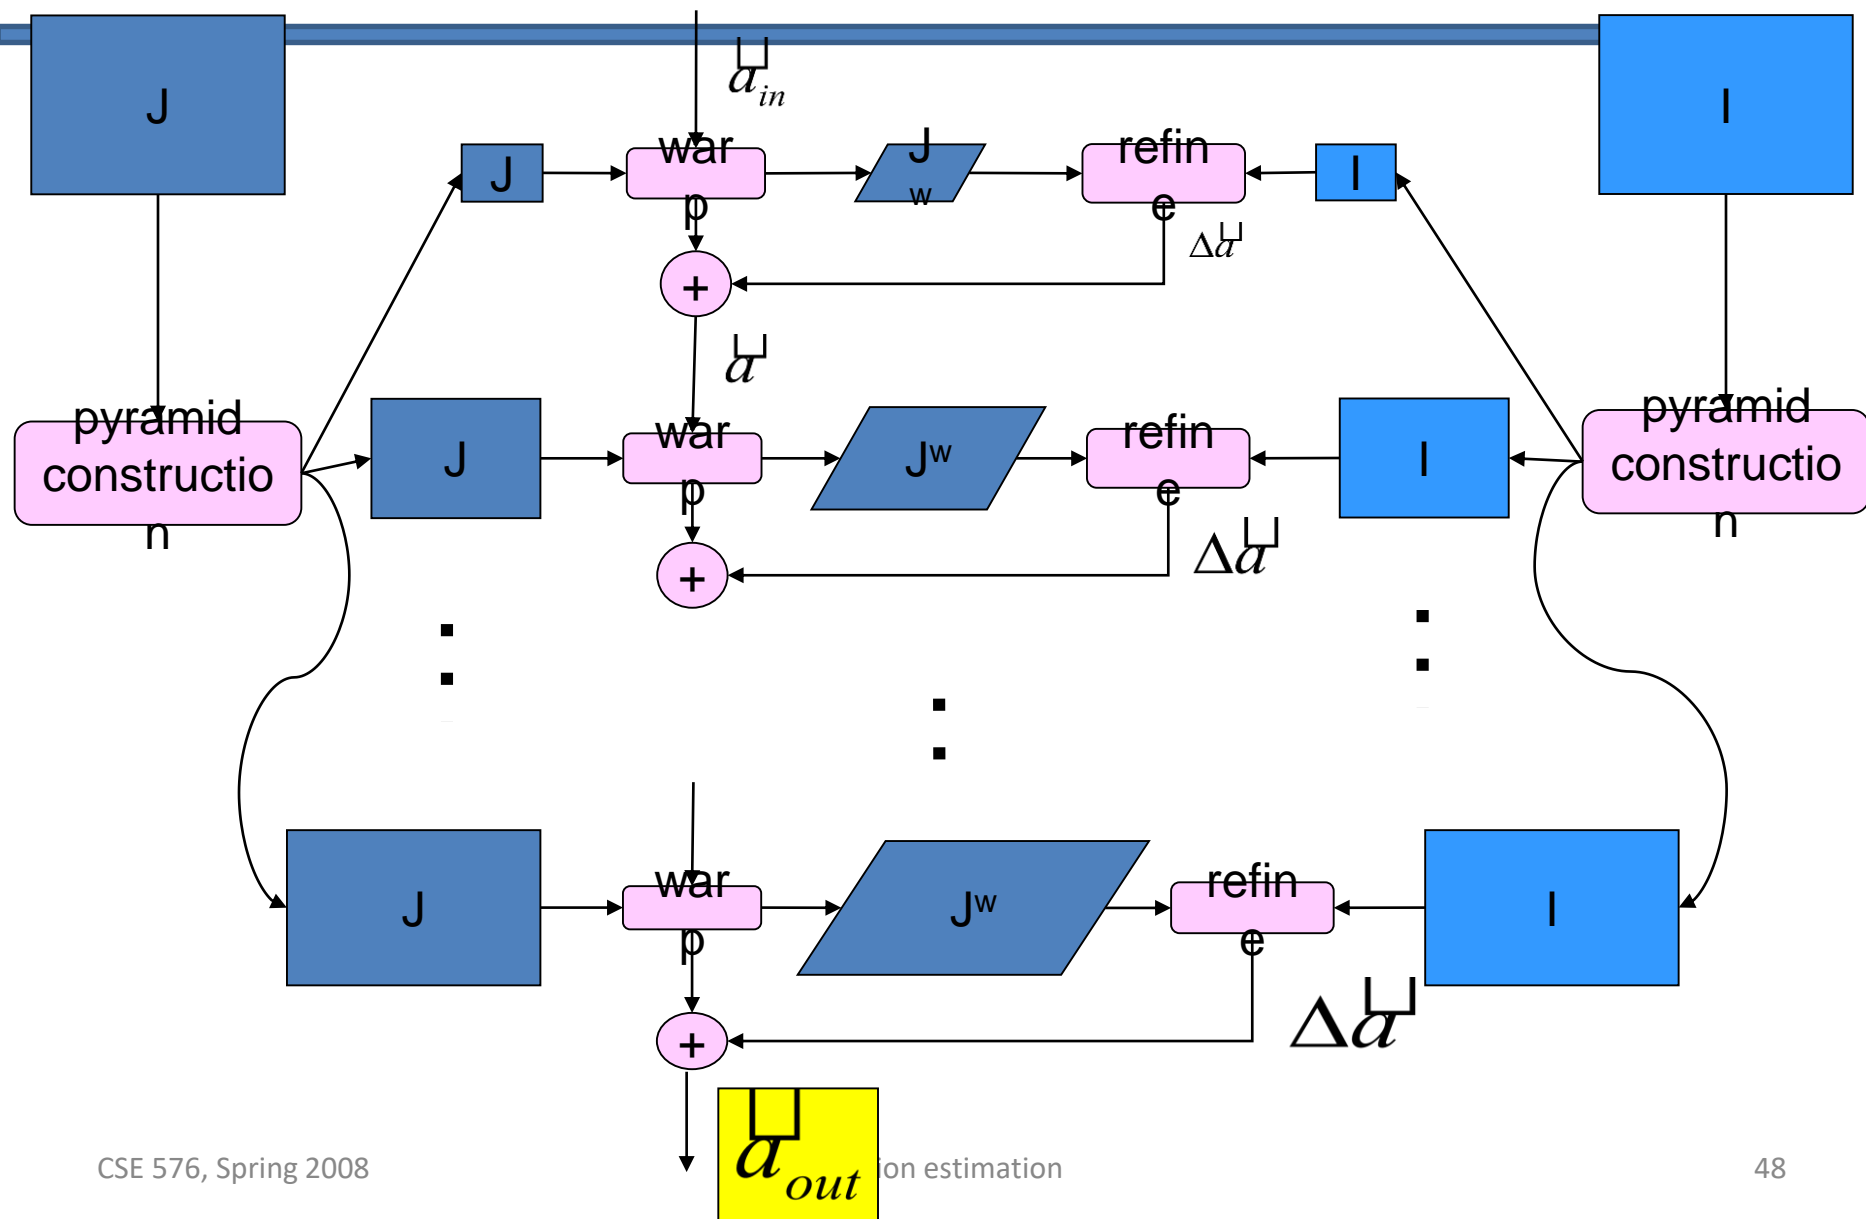
Coarse-to-Fine Estimation



Coarse-to-fine optical flow estimation




Coarse-to-Fine Estimation




Multi-resolution Lucas Kanade Algorithm

- Compute ‘simple’ LK at highest level
- At level i
 - Take flow u_{i-1}, v_{i-1} from level $i-1$
 - bilinear interpolate it to create u_i^*, v_i^* matrices of twice resolution for level i
 - multiply u_i^*, v_i^* by 2
 - compute f_t from a block displaced by $u_i^*(x,y), v_i^*(x,y)$
 - Apply LK to get $u_i'(x,y), v_i'(x,y)$ (the correction in flow)
 - Add corrections u_i', v_i' , *i.e.* $u_i = u_i^* + u_i'$, $v_i = v_i^* + v_i'$.


Optical Flow Results



Optical Flow Results




Lucas-Kanade with Pyramids



Feature tracking

- So far, we have only considered optical flow estimation in a pair of images
- If we have more than two images, we can compute the optical flow from each frame to the next
- Given a point in the first image, we can in principle reconstruct its path by simply “following the arrows”



Tracking challenges

- Ambiguity of optical flow
 - Need to find good features to track
- Large motions, changes in appearance, occlusions, disocclusions
 - Need mechanism for deleting, adding new features
- Drift – errors may accumulate over time
 - Need to know when to terminate a track


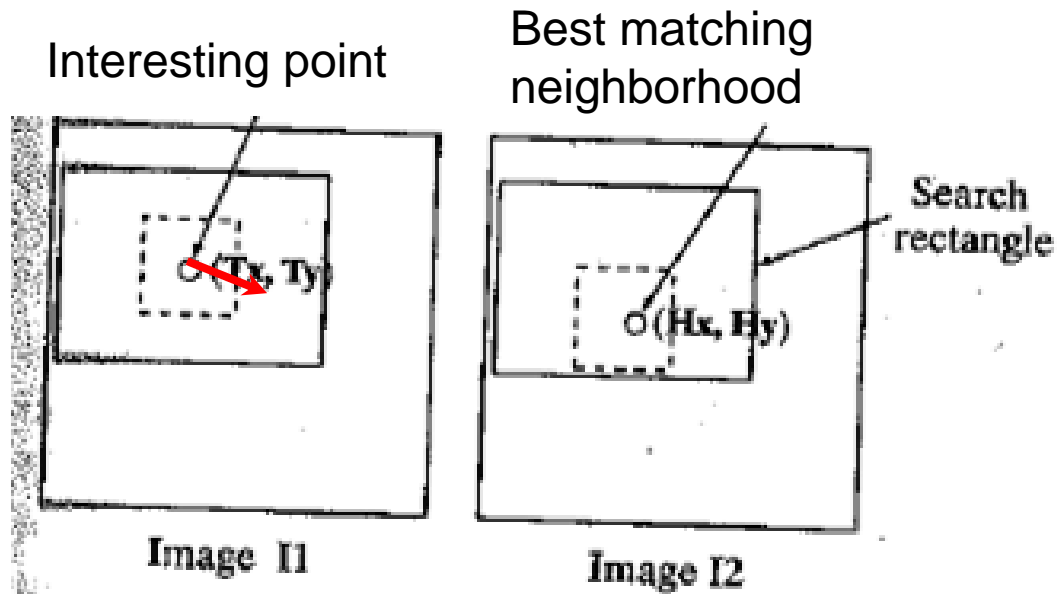
Tracking over many frames

- Select features in first frame
- For each frame:
 - Update positions of tracked features
 - Discrete search or Lucas-Kanade (or a combination of the two)
 - Terminate inconsistent tracks
 - Compute similarity with corresponding feature in the previous frame or in the first frame where it's visible
 - Find more features to track

What are good features to track?

- Harris corner detector
 - Can measure quality of features from just a single image
 - Automatically select candidate “templates”

Feature-based matching for motion



Shi-Tomasi feature tracker

1. Find good features (min eigenvalue of 2×2 Hessian)
2. Use Lucas-Kanade to track with pure translation
3. Use affine registration with first feature patch
4. Terminate tracks whose dissimilarity gets too large
5. Start new tracks when needed

A Camera Mouse

- Video interface: use feature tracking as mouse replacement
 - User clicks on the feature to be tracked
 - Take the 15x15 pixel square of the feature
 - In the next image do a search to find the 15x15 region with the highest correlation
 - Move the mouse pointer accordingly
 - Repeat in the background every 1/30th of a second



Implementation issues

- Window size
 - Small window more sensitive to noise and may miss larger motions (without pyramid)
 - Large window more likely to cross an occlusion boundary (and it's slower)
 - 15x15 to 31x31 seems typical
- Weighting the window
 - Common to apply weights so that center matters more (e.g., with Gaussian)

Optical Flow: Iterative Estimation

- Some Implementation Issues:
 - Warping is not easy (ensure that errors in warping are smaller than the estimate refinement)
 - Warp one image, take derivatives of the other so you don't need to re-compute the gradient after each iteration.
 - Often useful to low-pass filter the images before motion estimation (for better derivative estimation, and linear approximations to image intensity)

Summary

- Motion field: 3d motions projected to 2d images; dependency on depth
- Solving for motion with
 - sparse feature matches
 - dense optical flow
- Optical flow – key ideas
 - Brightness constancy assumption
 - Aperture problem
 - Solution with spatial coherence assumption
 - Coarse-to-fine registration